

Lecture 3: Data Manipulation

October 10, 2017

Contents

- `apply` family functions
- Intro to `dplyr` package
- Data manipulation functions
- Chaining operations
- Joining & appending datasets
- Extras

apply family

apply, lapply, sapply functions

- The `apply` family functions, are **functions which manipulate slices of data stored as matrices, arrays, lists and data-frames in a repetitive way.**
- These functions **avoid the explicit use of loops**, and might be more **computationally efficient**, depending on how big a dataset is. For more details on runtimes see this [link](#).
- `apply` allow you to perform operations with **very few lines of code.**
- The family comprises: **`apply`, `lapply`, `sapply`, `vapply`, `mapply`, `rapply`, and `tapply`**. The difference lies in the structure of input data and the desired format of the output).

apply function

apply operates on arrays/matrices.

In the example below we obtain column sums of matrix X.

```
(X <- matrix(sample(30), nrow = 5, ncol = 6))
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6]  
## [1, ] 23  25  29   1  13   8  
## [2, ]  3   4   6  12  22   2  
## [3, ]  5  21  26  17  19  14  
## [4, ]  7  10  15  27  16  28  
## [5, ] 24  18  20   9  30  11
```

```
apply(X, MARGIN = 2, FUN = sum)
```

```
## [1] 62 78 96 66 100 63
```

Note: that in a matrix `MARGIN = 1` indicates rows and `MARGIN = 2` indicates columns.

apply function

- `apply` can be used with **user-defined functions**:

```
print(X)
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6]  
## [1, ]  23  25  29   1  13   8  
## [2, ]   3   4   6  12  22   2  
## [3, ]   5  21  26  17  19  14  
## [4, ]   7  10  15  27  16  28  
## [5, ]  24  18  20   9  30  11
```

```
# number entries < 15  
apply(X, 2, function(x) 10*x + 2)
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6]  
## [1, ] 232 252 292  12 132  82  
## [2, ]  32  42  62 122 222  22  
## [3, ]  52 212 262 172 192 142  
## [4, ]  72 102 152 272 162 282  
## [5, ] 242 182 202  92 302 112
```

- The function can be defined outside `apply()`,

```
logColMeans <- function(x, eps = NULL)  
  if (!is.null(eps)) x <- x + eps  
  return(mean(x))  
}  
apply(X, 2, logColMeans)
```

```
## [1] 12.4 15.6 19.2 13.2 20.0 12.6
```

```
apply(X, 2, logColMeans, eps = 0.1)
```

```
## [1] 12.5 15.7 19.3 13.3 20.1 12.7
```

lapply/sapply functions

- `lapply()` is used to **repeatedly apply a function to elements of a sequential object** such as a vector, list, or data-frame (applies to columns).
- The **output is a list** with the same number of elements as the input object.
- `sapply` is the same as `lapply` but **returns a “simplified” output**.
- user-defined functions can be used with `sapply/lapply`

```
# lapply returns a list  
lapply(1:3, function(x) x^2)
```

```
## [[1]]  
## [1] 1  
##  
## [[2]]  
## [1] 4  
##  
## [[3]]  
## [1] 9
```

```
# which you can 'simplify' with unlist()  
unlist(lapply(1:3, function(x) x^2))
```

```
## [1] 1 4 9
```

```
# Or you could use sapply() instead  
sapply(1:3, function(x) x^2)
```

```
## [1] 1 4 9
```

mapply functions

- mapp`l`y stands for ‘multivariate’ apply. It **applies a function to a multiple list or multiple vectors as arguments**.
- The goal is to vectorize arguments to a function which usually does not accept vectors as arguments.

```
# function word() returns a string of character C repeated k times.
word <- function(C,k) paste(rep.int(C,k), collapse='')
mapply(word, LETTERS[1:6], 6:1, SIMPLIFY = FALSE)
```

```
## $A
## [1] "AAAAAA"
##
## $B
## [1] "BBBBB"
##
## $C
## [1] "CCCC"
##
## $D
## [1] "DDD"
##
## $E
## [1] "EE"
##
## $F
## [1] "F"
```


Exercise 1

- Go to the “Lec3_Exercises.Rmd” file, which can be downloaded from the class website under the Lecture tab.
- Complete Exercise 1.

dpLyr package

dplyr

- Introduces a **grammar of data manipulation**.
- **Code-efficient** for data exploration and transformation.
- **Fast on data frames** (written in C++): has speed of C and ease of R.
- **Intuitive to write and easy to read**, esp. when using the *chaining* syntax.
- You should use dplyr even as a beginner R user, and [here is why](#).

```
# To install dplyr with latest updates
install.packages("devtools")
devtools::install_github("tidyverse/dplyr")

# Or you could use CRAN
install.packages("dplyr")
```

tibbles

Tibbles are a modern take on data frames. They keep the features that have stood the test of time, and drop the features that used to be convenient but are now frustrating (i.e. converting character vectors to factors).

- A tibble, `tbl`, is a **wrapper for a data frame that prints nicely**.
- The print method for `tbl` shows only the **first 10 rows, and all the columns that fit on screen**.
- Each column is also reported together its type.
- Tibbles (and `dplyr`) **do NOT preserve the row names**.
- **Subsetting `tbl` is stricter** than subsetting `data.frames`, and **ALWAYS** returns objects with expected class, i.e. with a single `[` you get back a tibble, and with `double [[` you get a vector.

Movie industry dataset

`movies.csv` contains information on last three decades of movies. The data has been scraped from the IMDb website and can be accessed from a [github repo](#).

```
url <- "https://raw.githubusercontent.com/Juanets/movie-stats/master/movies.csv"
movies.df <- read.csv(url)
rownames(movies.df) <- paste0("M", 1:nrow(movies.df))
dim(movies.df)
```

```
## [1] 6820 15
```

```
colnames(movies.df)
```

```
## [1] "budget" "company" "country" "director" "genre" "gross"
## [7] "name" "rating" "released" "runtime" "score" "star"
## [13] "votes" "writer" "year"
```

head(movies.df)

##	budget		company		country	director		
##	M1	8000000	Columbia Pictures Corporation		USA	Rob Reiner		
##	M2	6000000	Paramount Pictures		USA	John Hughes		
##	M3	15000000	Paramount Pictures		USA	Tony Scott		
##	M4	18500000	Twentieth Century	Fox Film Corporation	USA	James Cameron		
##	M5	9000000	Walt Disney Pictures		USA	Randal Kleiser		
##	M6	6000000	Hemdale		UK	Oliver Stone		
##	genre	gross	name		rating	released	runtime	score
##	M1	Adventure	52287414	Stand by Me	R	1986-08-22	89	8.1
##	M2	Comedy	70136369	Ferris Bueller's Day Off	PG-13	1986-06-11	103	7.5
##	M3	Action	179800601	Top Gun	PG	1986-05-16	110	6.8
##	M4	Action	85160248	Aliens	R	1986-07-18	137	8.4
##	M5	Adventure	18564613	Flight of the Navigator	PG	1986-08-01	90	6.5
##	M6	Drama	138530565	Platoon	R	1987-02-06	120	8.2
##	star	votes	writer		year			
##	M1	Wil Wheaton	299174	Stephen King	1986			
##	M2	Matthew Broderick	264740	John Hughes	1986			
##	M3	Tom Cruise	236909	Jim Cash	1986			
##	M4	Sigourney Weaver	540152	James Cameron	1986			
##	M5	Joey Cramer	36636	Mark H. Baker	1986			
##	M6	Charlie Sheen	317585	Oliver Stone	1986			

Convert to tibble

```
library(dplyr)

# convert to tibble
movies <- tbl_df(movies.df)
class(movies)
```

```
## [1] "tbl_df"      "tbl"        "data.frame"
```

```
# printing only shows 10 rows and as many columns as can fit on your screen
movies
```

```
## # A tibble: 6,820 x 15
##   budget company country director
##   *    <dbl>          <fctr>   <fctr>      <fctr>
## 1  8000000 Columbia Pictures Corporation USA Rob Reiner
## 2  6000000 Paramount Pictures USA John Hughes
## 3 15000000 Paramount Pictures USA Tony Scott
## 4 18500000 Twentieth Century Fox Film Corporation USA James Cameron
## 5  9000000 Walt Disney Pictures USA Randal Kleiser
## 6  6000000 Hemdale UK Oliver Stone
## 7 25000000 Henson Associates (HA) UK Jim Henson
## 8  6000000 De Laurentiis Entertainment Group (DEG) USA David Lynch
## 9  9000000 Paramount Pictures USA Howard Deutch
## 10 15000000 SLM Production Group USA David Cronenberg
## # ... with 6,810 more rows, and 11 more variables: genre <fctr>, gross <dbl>,
## #   name <fctr>, rating <fctr>, released <fctr>, runtime <int>, score <dbl>,
## #   star <fctr>, votes <int>, writer <fctr>, year <int>
```

Note: Tibbles and dplyr do NOT preserve the row names, you need to create an “id” column instead.

```
str(movies)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':    6820 obs. of  15 variables:
## $ budget   : num  80000000 60000000 150000000 185000000 90000000 60000000 250000000 (
## $ company  : Factor w/ 2179 levels "101st Street Films",...: 663 1683 1683 2068
## $ country  : Factor w/ 57 levels "Argentina","Aruba",...: 56 56 56 56 56 54 54
## $ director : Factor w/ 2759 levels "Aamir Khan","Aaron Blaise",...: 2246 1295 2
## $ genre     : Factor w/ 17 levels "Action","Adventure",...: 2 5 1 1 2 7 2 7 5 7
## $ gross     : num  5.23e+07 7.01e+07 1.80e+08 8.52e+07 1.86e+07 ...
## $ name      : Factor w/ 6731 levels "10,000 BC","101 Dalmatians",...: 4667 1819
## $ rating    : Factor w/ 13 levels "B","B15","G",...: 9 8 7 9 7 9 7 9 8 9 ...
## $ released  : Factor w/ 2403 levels "1986-01-10","1986-01-17",...: 40 28 24 34 3
## $ runtime   : int    89 103 110 137 90 120 101 120 96 96 ...
## $ score     : num    8.1 7.8 6.9 8.4 6.9 8.1 7.4 7.8 6.8 7.5 ...
## $ star      : Factor w/ 2504 levels "50 Cent","Aaliyah",...: 2474 1605 2348 2195
## $ votes     : int    299174 264740 236909 540152 36636 317585 102879 146768 60565
## $ writer    : Factor w/ 4199 levels "A.A. Milne","Aaron Guzikowski",...: 3723 19
## $ year      : int    1986 1986 1986 1986 1986 1986 1986 1986 1986 1986 ...
```


Data manipulation functions

dp̃yr verbs

The most commonly used dp̃yr functions (or basic verbs) are:

- `filter()`: keep rows matching criteria,
- `select()`: pick columns by name,
- `arrange()`: reorder rows,
- `mutate()`: add new variables,
- `summarise()`: reduce variables to values

Operations performed with the above functions can be done using **base R functions**, but they would be less computationally efficient, and require writing more lines of (ugly) code.

Learn about dp̃yr from the [tutorial](#) written by its creator, Hadley Wickham

Structure of `dplyr` functions

- the first argument is a data frame
- subsequent argument specify what to do
- always return a data frame

filter(): keep rows matching criteria

```
# base R approach to find all comedies by Woody Allen  
movies[movies$genre == "Comedy" & movies$director == "Woody Allen", ]
```

```
# dplyr approach  
# note: both comma or ampersand represent AND condition  
filter(movies, genre == "Comedy", director == "Woody Allen")
```

```
## # A tibble: 27 x 15  
##       budget company country director  
##       <dbl>      <fctr>   <fctr>   <fctr>  
## 1  64000000 Orion Pictures USA Woody Allen  
## 2 160000000 Orion Pictures USA Woody Allen  
## 3 190000000 Jack Rollins & Charles H. Joffe Productions USA Woody Allen  
## 4 150000000 Touchstone Pictures USA Woody Allen  
## 5 120000000 Orion Pictures USA Woody Allen  
## 6 140000000 Orion Pictures USA Woody Allen  
## 7 200000000 TriStar Pictures USA Woody Allen  
## 8 135000000 TriStar Pictures USA Woody Allen  
## 9 200000000 Miramax USA Woody Allen  
## 10 150000000 Sweetland Films USA Woody Allen  
## # ... with 17 more rows, and 11 more variables: genre <fctr>, gross <dbl>,  
## #   name <fctr>, rating <fctr>, released <fctr>, runtime <int>, score <dbl>,  
## #   star <fctr>, votes <int>, writer <fctr>, year <int>
```

```
# use pipe for OR condition
filter(movies, country == "Greece" | country == "Chile")
```

```
## # A tibble: 9 x 15
##   budget          company country      director      genre
##   <dbl>          <fctr>   <fctr>      <fctr>      <fctr>
## 1 0.0e+00      Paradis Films Greece Theodoros Angelopoulos Drama
## 2 0.0e+00      Boo Productions Greece Yorgos Lanthimos Drama
## 3 0.0e+00      Haos Film Greece Athina Rachel Tsangari Drama
## 4 0.0e+00      Haos Film Greece Yorgos Lanthimos Drama
## 5 0.0e+00 Participant Media Chile "Pablo Larra\xeddn" Drama
## 6 0.0e+00      Film4 Greece Yorgos Lanthimos Comedy
## 7 2.6e+07      Alcon Entertainment Chile Patricia Ríígen Biography
## 8 9.0e+06 Fox Searchlight Pictures Chile "Pablo Larra\xeddn" Biography
## 9 0.0e+00      AZ Films Chile "Pablo Larra\xeddn" Biography
## # ... with 10 more variables: gross <dbl>, name <fctr>, rating <fctr>,
## #   released <fctr>, runtime <int>, score <dbl>, star <fctr>, votes <int>,
## #   writer <fctr>, year <int>
```

```
# you can also use %in% operator
print(filter(movies, country %in% c("Argentina", "Colombia", "Chile")), n = Inf)
```

```
## # A tibble: 19 x 15
##   budget company country
##   <dbl> <fctr> <fctr>
## 1 0.0e+00 Cinequanon Argentina
## 2 0.0e+00 "GEA Cinematogr\xe1fica" Argentina
## 3 0.0e+00 Not specified Argentina
## 4 0.0e+00 Aleph Producciones S.A. Argentina
## 5 1.5e+06 FX Sound Argentina
## 6 0.0e+00 Instituto Nacional de Cine y Artes Audiovisuales (INCAA) Argentina
## 7 0.0e+00 4k Films Argentina
## 8 0.0e+00 FilmFour Argentina
## 9 3.0e+06 HBO Films Colombia
## 10 0.0e+00 Cinefarm Argentina
## 11 0.0e+00 Aura Films Argentina
## 12 0.0e+00 Historias Cinematograficas Cinemania Argentina
## 13 2.0e+06 Tornasol Films Argentina
## 14 0.0e+00 Participant Media Chile
## 15 3.3e+06 Corner Producciones Argentina
## 16 2.6e+07 Alcon Entertainment Chile
## 17 1.4e+06 Buffalo Films Colombia
## 18 9.0e+06 Fox Searchlight Pictures Chile
## 19 0.0e+00 AZ Films Chile
## # ... with 12 more variables: director <fctr>, genre <fctr>, gross <dbl>,
## #   name <fctr>, rating <fctr>, released <fctr>, runtime <int>, score <dbl>,
## #   star <fctr>, votes <int>, writer <fctr>, year <int>
```

select(): pick columns by name

```
# base R approach to select columns
movies[, c("name", "year", "genre")]
```

```
# dplyr approach
movies.sub <- select(movies, name, country, year, genre)
movies.sub
```

```
## # A tibble: 6,820 x 4
##       name country year genre
##   *      <fctr>  <fctr> <int> <fctr>
## 1      Stand by Me    USA  1986 Adventure
## 2 Ferris Bueller's Day Off    USA  1986 Comedy
## 3          Top Gun    USA  1986 Action
## 4          Aliens    USA  1986 Action
## 5 Flight of the Navigator    USA  1986 Adventure
## 6          Platoon    UK  1986 Drama
## 7          Labyrinth    UK  1986 Adventure
## 8          Blue Velvet    USA  1986 Drama
## 9          Pretty in Pink    USA  1986 Comedy
## 10         The Fly    USA  1986 Drama
## # ... with 6,810 more rows
```

```
# use colon to select multiple contiguous columns,  
select(movies, name, genre:score)
```

```
## # A tibble: 6,820 x 7  
##           name      genre      gross rating  released runtime score  
##    *      <fctr>    <fctr>    <dbl> <fctr>    <fctr>    <int> <dbl>  
##  1      Stand by Me Adventure  52287414      R 1986-08-22      89      8.  
##  2 Ferris Bueller's Day Off Comedy  70136369 PG-13 1986-06-11     103      7.  
##  3          Top Gun   Action 179800601      PG 1986-05-16     110      6.  
##  4          Aliens   Action  85160248      R 1986-07-18     137      8.  
##  5 Flight of the Navigator Adventure  18564613      PG 1986-08-01      90      6.  
##  6          Platoon   Drama 138530565      R 1987-02-06     120      8.  
##  7          Labyrinth Adventure  12729917      PG 1986-06-27     101      7.  
##  8          Blue Velvet Drama    8551228      R 1986-10-23     120      7.  
##  9          Pretty in Pink Comedy  40471663 PG-13 1986-02-28      96      6.  
## 10          The Fly   Drama  40456565      R 1986-08-15      96      7.  
## # ... with 6,810 more rows
```


select() helpers

You can use the following functions to help select the columns:

- starts_with()
- ends_with()
- contains()
- matches() (matches a regular expression)

```
select(movies, starts_with("r"))
```

```
## # A tibble: 6,820 x 3
##   rating    released runtime
##   * <fctr>      <fctr>    <int>
## 1      R 1986-08-22      89
## 2 PG-13 1986-06-11     103
## 3      PG 1986-05-16     110
## 4      R 1986-07-18     137
## 5      PG 1986-08-01      90
## 6      R 1987-02-06     120
## 7      PG 1986-06-27     101
## 8      R 1986-10-23     120
## 9 PG-13 1986-02-28      96
## 10     R 1986-08-15      96
## # ... with 6,810 more rows
```

```
select(movies, ends_with("e"))
```

```
## # A tibble: 6,820 x 4
##   genre      name runtime score
##   <fctr>    <fctr>   <int> <dbl>
## 1 Adventure Stand by Me      89    8.1
## 2 Comedy   Ferris Bueller's Day Off 103    7.8
## 3 Action    Top Gun      110    6.9
## 4 Action    Aliens       137    8.4
## 5 Adventure Flight of the Navigator    90    6.9
## 6 Drama     Platoon      120    8.1
## 7 Adventure Labyrinth    101    7.4
## 8 Drama     Blue Velvet   120    7.8
## 9 Comedy   Pretty in Pink   96    6.8
## 10 Drama    The Fly        96    7.5
## # ... with 6,810 more rows
```

```
select(movies, contains("re"))
```

```
## # A tibble: 6,820 x 4
##   director      genre  released score
##   <fctr>    <fctr>    <fctr> <dbl>
## 1 Rob Reiner Adventure 1986-08-22    8.1
## 2 John Hughes Comedy  1986-06-11    7.8
## 3 Tony Scott  Action  1986-05-16    6.9
## 4 James Cameron Action  1986-07-18    8.4
## 5 Randal Kleiser Adventure 1986-08-01    6.9
## 6 Oliver Stone Drama    1987-02-06    8.1
## 7 Jim Henson Adventure 1986-06-27    7.4
## 8 David Lynch Drama    1986-10-23    7.8
## 9 Howard Deutch Comedy  1986-02-28    6.8
## 10 David Cronenberg Drama    1986-08-15    7.5
## # ... with 6,810 more rows
```

Dropping columns

```
# remove budget and company columns
print(select(movies, -budget, -company), n = 6)
```

```
## # A tibble: 6,820 x 13
##   country      director      genre      gross      name rating
##   <fctr>      <fctr>      <fctr>      <dbl>      <fctr> <fctr>
## 1      USA      Rob Reiner Adventure  52287414      Stand by Me      R
## 2      USA      John Hughes  Comedy   70136369 Ferris Bueller's Day Off PG-13
## 3      USA      Tony Scott   Action  179800601      Top Gun      PG
## 4      USA      James Cameron Action   85160248      Aliens      R
## 5      USA      Randal Kleiser Adventure  18564613 Flight of the Navigator PG
## 6      UK      Oliver Stone  Drama   138530565      Platoon      R
## # ... with 6,814 more rows, and 7 more variables: released <fctr>,
## #   runtime <int>, score <dbl>, star <fctr>, votes <int>, writer <fctr>,
## #   year <int>
```

```
# Selecting and renaming in one
print(select(movies, name, gross_revenue = gross), n = 6)
```

```
## # A tibble: 6,820 x 2
##   name      gross_revenue
##   <fctr>      <dbl>
## 1      Stand by Me      52287414
## 2 Ferris Bueller's Day Off  70136369
## 3      Top Gun      179800601
## 4      Aliens      85160248
## 5 Flight of the Navigator  18564613
## 6      Platoon      138530565
## # ... with 6,814 more rows
```

arrange(): reorder rows

```
# dplyr approach
print(arrange(movies.sub, name), n = 6)
```

```
## # A tibble: 6,820 x 4
##       name country year genre
##       <fctr> <fctr> <int> <fctr>
## 1 10,000 BC USA 2008 Action
## 2 101 Dalmatians USA 1996 Adventure
## 3 102 Dalmatians USA 2000 Adventure
## 4 10 Cloverfield Lane USA 2016 Drama
## 5 10 Things I Hate About You USA 1999 Comedy
## 6 10 Years USA 2011 Comedy
## # ... with 6,814 more rows
```

```
# use `desc` for descending
print(arrange(movies.sub, desc(year), n = 6))
```

```
## # A tibble: 6,820 x 4
##       name country year genre
##       <fctr> <fctr> <int> <fctr>
## 1 The Bad Batch USA 2016 Drama
## 2 Assassin's Creed USA 2016 Action
## 3 La La Land USA 2016 Comedy
## 4 Suicide Squad USA 2016 Action
## 5 Sing USA 2016 Animation
## 6 Fantastic Beasts and Where to Find Them UK 2016 Adventure
## # ... with 6,814 more rows
```

mutate(): add new variables

```
# base R approach to create a new variable 'profit'  
movies$profit <- movies$gross - movies$budget
```

```
# dplyr approach  
movies <- mutate(movies, profit = gross - budget)  
select(movies, name, gross, budget, profit)
```

```
## # A tibble: 6,820 x 4  
##           name      gross    budget    profit  
##           <fctr>    <dbl>    <dbl>    <dbl>  
## 1      Stand by Me  52287414  80000000  44287414  
## 2 Ferris Bueller's Day Off  70136369  60000000  64136369  
## 3        Top Gun  179800601  150000000  164800601  
## 4         Aliens   85160248  185000000  66660248  
## 5 Flight of the Navigator  18564613   90000000   9564613  
## 6         Platoon  138530565   60000000  132530565  
## 7       Labyrinth  12729917  250000000 -12270083  
## 8      Blue Velvet   8551228   60000000   2551228  
## 9    Pretty in Pink  40471663   90000000  31471663  
## 10      The Fly    40456565  150000000  25456565  
## # ... with 6,810 more rows
```

Generating multiple new variables

```
movies <- mutate(movies,  
  profit = gross - budget,  
  gross_in_mil = gross/10^6,  
  budget_in_mil = budget/10^6,  
  profit_in_mil = profit/10^6)  
select(movies, name, contains("_in_mil"))
```

```
## # A tibble: 6,820 x 4  
##           name gross_in_mil budget_in_mil profit_in_mil  
##           <fctr>         <dbl>         <dbl>         <dbl>  
## 1      Stand by Me      52.287414          8.0         44.287414  
## 2 Ferris Bueller's Day Off 70.136369          6.0         64.136369  
## 3        Top Gun    179.800601         15.0        164.800601  
## 4        Aliens     85.160248         18.5         66.660248  
## 5 Flight of the Navigator  18.564613          9.0          9.564613  
## 6        Platoon   138.530565          6.0       132.530565  
## 7      Labyrinth   12.729917         25.0        -12.270083  
## 8      Blue Velvet    8.551228          6.0          2.551228  
## 9    Pretty in Pink   40.471663          9.0         31.471663  
## 10     The Fly      40.456565         15.0         25.456565  
## # ... with 6,810 more rows
```

summarise(): reduce variables to values

- summarise() can be used to aggregate data or to compute a summarizing value of interest.
- summarise() is primarily useful data previously grouped by one or more variables using group_by()

```
# Total gross revenue for all movies in 1969
summarise(movies,
  tot_gross_in_bil = sum(gross)/
```

```
## # A tibble: 1 x 1
##   tot_gross_in_bil
##   <dbl>
## 1      228.4552
```

```
by_genre <- group_by(movies, genre)
summarise(by_genre, total = sum(gross))
```

```
## # A tibble: 17 x 2
##   genre      total
##   <fctr>    <dbl>
## 1 Action  74.792634664
## 2 Adventure 20.895092830
## 3 Animation 25.342203262
## 4 Biography  8.617526666
## 5 Comedy  53.543423603
## 6 Crime  10.217836159
## 7 Drama  25.204622256
## 8 Family  0.118110208
## 9 Fantasy  0.644653115
## 10 Horror  7.117846856
## 11 Musical  0.008094993
## 12 Mystery  1.379951556
## 13 Romance  0.145764206
## 14 Sci-Fi  0.307801697
## 15 Thriller  0.099600218
## 16 War  0.001509775
## 17 Western  0.018519094
```

Useful summary functions


- `min(x), median(x), max(x), quantile(x, p)`
- `n(), n_distinct(), sum(x), mean(x)`
- `sum(x > 10), mean(x > 0)`
- `sd(x), var(x)`

Chaining operations

Pipe operator: %>%

- The %>% operator was introduced in `magrittr` package to **pipe values forward into an expression or function call**.
- In the pipe notation we have something like `x %>% f(y)`, rather than `f(x, y)`
- This is similar to the **Unix pipes**: | used to **send the output of one program to another program for further processing**.

```
rnorm(500) %>%  
  matrix(ncol = 50) %>%  
  colSums() %>%  
  plot(x = 1:50)
```



Chaining operations

- Pipe operators used together with `dplyr` functions make a large difference as they semantically change your code in a way that **makes it more intuitive to both read and write**.
- The pipes allow users to chain operators which reflects the **sequential nature of data-processing tasks**.
- **Chaining increases readability** significantly when there are many commands
- `%>%` operator is automatically imported into `dplyr`

Find movies from USA produced after 2010. Group by genre and compute the group mean gross revenue in million dollars. Then print the genre mean 'gross' revenue arranged in a descending order:

```
# nesting
arrange(
  summarise(
    group_by(
      filter(movies,
              year > 2010, country == "US
            ),
            genre
          ),
    mean_gross = mean(gross)/10^6
  ),
  mean_gross
)
```

```
## # A tibble: 13 x 2
##   genre mean_gross
##   <fctr>    <dbl>
## 1 Thriller    0.016458
## 2 Drama     23.252067
## 3 Horror     27.714873
## 4 Sci-Fi     29.169033
## 5 Fantasy     30.657051
## 6 Crime      32.128370
## 7 Comedy     35.152773
## 8 Biography  40.552788
## 9 Mystery    49.457150
## 10 Romance   62.495645
## 11 Adventure  81.227901
## 12 Action    97.285930
## 13 Animation 151.960486
```

```
# chaining
movies %>%
  filter(year > 2010, country == "USA")
  group_by(genre) %>%
  summarise(mean_gross = mean(gross)/10^6)
  arrange(mean_gross)
```

```
## # A tibble: 13 x 2
##   genre mean_gross
##   <fctr>    <dbl>
## 1 Thriller    0.016458
## 2 Drama     23.252067
## 3 Horror     27.714873
## 4 Sci-Fi     29.169033
## 5 Fantasy     30.657051
## 6 Crime      32.128370
## 7 Comedy     35.152773
## 8 Biography  40.552788
## 9 Mystery    49.457150
## 10 Romance   62.495645
## 11 Adventure  81.227901
## 12 Action    97.285930
## 13 Animation 151.960486
```

Exercises 2

- Go to the “Lec3_Exercises.Rmd” file, which can be downloaded from the class website under the Lecture tab.
- Complete Exercise 2.

Joining & appending datasets

Merging two tables

```
x <- data.frame(
  item = c("cherries", "orange",
           "broccoli", "blueberries",
           "cabbage", "lemon", "spinach",
           "eggplants"),
  color = c("red", "orange", "green",
            "blue", "green", "yellow",
            "green", "purple"),
  stringsAsFactors = FALSE)

y <- data.frame(
  item = c("carrot", "cherries",
           "broccoli", "blueberries",
           "cabbage", "lemon", "tomato"),
  class = c("vegetable", "fruit",
            "vegetable", "fruit",
            "vegetable", "fruit", "fruit"),
  stringsAsFactors = FALSE)
```

item	color	item	class
cherries	red	carrot	vegetable
orange	orange	cherries	fruit
broccoli	green	broccoli	vegetable
blueberries	blue	blueberries	fruit
cabbage	green	cabbage	vegetable
lemon	yellow	lemon	fruit
spinach	green	tomato	fruit
eggplants	purple		

inner_join(x, y)

- Include only rows in both x and y

```
inner_join(x, y) # or x %>% inner_join(y)
```

```
## Joining, by = "item"
```

x		y	
item	color	item	class
cherries	red	carrot	vegetable
orange	orange	cherries	fruit
broccoli	green	broccoli	vegetable
blueberries	blue	blueberries	fruit
cabbage	green	cabbage	vegetable
lemon	yellow	lemon	fruit
spinach	green	tomato	fruit
eggplants	purple		

inner_join(x, y)		
item	color	class
cherries	red	fruit
broccoli	green	vegetable
blueberries	blue	fruit
cabbage	green	vegetable
lemon	yellow	fruit

left_join(x, y):

- Include all of x, and matching rows of y

```
left_join(x, y) # or x %>% left_join(y)
```

```
## Joining, by = "item"
```

x		y	
item	color	item	class
cherries	red	carrot	vegetable
orange	orange	cherries	fruit
broccoli	green	broccoli	vegetable
blueberries	blue	blueberries	fruit
cabbage	green	cabbage	vegetable
lemon	yellow	lemon	fruit
spinach	green	tomato	fruit
eggplants	purple		

left_join(x, y)		
item	color	class
cherries	red	fruit
orange	orange	NA
broccoli	green	vege
blueberries	blue	fruit
cabbage	green	vege
lemon	yellow	fruit
spinach	green	NA
eggplants	purple	NA

semi_join(x, y):

- Include rows of x that match y

```
semi_join(x, y) # or x %>% semi_join(y)
```

```
## Joining, by = "item"
```

x		y	
item	color	item	class
cherries	red	carrot	vegetable
orange	orange	cherries	fruit
broccoli	green	broccoli	vegetable
blueberries	blue	blueberries	fruit
cabbage	green	cabbage	vegetable
lemon	yellow	lemon	fruit
spinach	green	tomato	fruit
eggplants	purple		

semi_join(x,	
item	co
cherries	re
broccoli	gr
blueberries	bl
cabbage	gr
lemon	ye

anti_join(x, y):

- Include rows of x that don't match y

```
anti_join(x, y) # or x %>% anti_join(y)
```

```
## Joining, by = "item"
```

x		y	
item	color	item	class
cherries	red	carrot	vegetable
orange	orange	cherries	fruit
broccoli	green	broccoli	vegetable
blueberries	blue	blueberries	fruit
cabbage	green	cabbage	vegetable
lemon	yellow	lemon	fruit
spinach	green	tomato	fruit
eggplants	purple		

anti_join(x,	
item	cc
orange	or
spinach	gr
eggplants	pu

Extra functions

Renaming

```
movies.sub <- movies %>% select(name, director, year, score, gross)
print(movies.sub, n = 3)
```

```
## # A tibble: 6,820 x 5
##           name      director  year score    gross
##           <fctr>    <fctr> <int> <dbl>    <dbl>
## 1      Stand by Me  Rob Reiner  1986   8.1  52287414
## 2 Ferris Bueller's Day Off John Hughes  1986   7.8  70136369
## 3          Top Gun  Tony Scott  1986   6.9 179800601
## # ... with 6,817 more rows
```

```
# Renaming variable
movies.sub %>% rename(gross_revenue = gross)
```

```
## # A tibble: 6,820 x 5
##           name      director  year score gross_revenue
##           <fctr>    <fctr> <int> <dbl>    <dbl>
## 1      Stand by Me  Rob Reiner  1986   8.1    52287414
## 2 Ferris Bueller's Day Off John Hughes  1986   7.8    70136369
## 3          Top Gun  Tony Scott  1986   6.9   179800601
## 4          Aliens  James Cameron  1986   8.4    85160248
## 5 Flight of the Navigator Randal Kleiser  1986   6.9    18564613
## 6          Platoon  Oliver Stone  1986   8.1   138530565
## 7          Labyrinth  Jim Henson  1986   7.4    12729917
## 8          Blue Velvet  David Lynch  1986   7.8     8551228
## 9      Pretty in Pink  Howard Deutch  1986   6.8    40471663
## 10         The Fly  David Cronenberg  1986   7.5    40456565
## # ... with 6,810 more rows
```

Distinct values

```
# Unique values  
movies %>% distinct(rating)
```

```
## # A tibble: 13 x 1  
##           rating  
##       <fctr>  
## 1             R  
## 2          PG-13  
## 3            PG  
## 4        UNRATED  
## 5 Not specified  
## 6             G  
## 7          NC-17  
## 8        NOT RATED  
## 9          TV-PG  
## 10         TV-MA  
## 11            B  
## 12          B15  
## 13         TV-14
```

Note that can take on multiple variables, and would return distinct variable combinations.

Group counts

Use a `tally()` function to generate a group frequency table:

```
movies %>% group_by(genre) %>% tally()
```

```
## # A tibble: 17 x 2
##   genre      n
##   <fctr> <int>
## 1 Action  1331
## 2 Adventure  392
## 3 Animation  277
## 4 Biography  359
## 5 Comedy  2080
## 6 Crime    522
## 7 Drama   1444
## 8 Family    14
## 9 Fantasy    32
## 10 Horror   277
## 11 Musical     4
## 12 Mystery    38
## 13 Romance    15
## 14 Sci-Fi     13
## 15 Thriller   18
## 16 War         2
## 17 Western     2
```

Window Functions

- Aggregation functions such as `mean()`, `n()` return 1 value per group.
- **Window functions return multiple values per group.** Examples include: ranking and ordering functions (like `min_rank`, `top_n()`), offset functions (`lead` and `lag`), and cumulative aggregates (like `cummean`).

```
# rewrite more simply with the `top_n` function
movies %>%
  select(name, genre, year, score) %>%
  group_by(genre) %>%
  top_n(2, wt = score) %>%           # if 'wt' argument in top_n() is not specified
  arrange(genre, year, score)       # the last variable is taken for ordering
```

```
## # A tibble: 35 x 4
## # Groups:   genre [17]
##           name      genre  year  score
##           <fctr>   <fctr> <int> <dbl>
## 1 The Dark Knight Action   2008   9.0
## 2 Inception        Action   2010   8.8
## 3 The Lord of the Rings: The Fellowship of the Ring Adventure 2001   8.8
## 4 The Lord of the Rings: The Return of the King Adventure 2003   8.9
## 5 The Lion King     Animation 1994   8.5
## 6 Spirited Away     Animation 2001   8.6
## 7 Your name         Animation 2016   8.5
## 8 Schindler's List Biography 1993   8.9
## 9 The Intouchables Biography 2011   8.6
## 10 Forrest Gump      Comedy   1994   8.8
## # ... with 25 more rows
```


Sampling functions

```
# randomly sample a fixed number of rows, without replacement
movies %>% sample_n(15)
```

```
## # A tibble: 15 x 19
##       budget          company country      director      genre
##       <dbl>          <fctr>   <fctr>      <fctr>      <fctr>
## 1 4.50e+06    Keystone Family Pictures    Canada    Robert Vince    Comed
## 2 0.00e+00 Raajkamal Films International    India      Kamal Haasan    Crim
## 3 1.75e+07      United Artists          UK        Terry George    Biograph
## 4 1.60e+07    New Line Cinema          USA      Jonathan Liebesman    Horro
## 5 0.00e+00      Warner Bros.          USA        John Gray    Adventur
## 6 2.00e+08      DreamWorks          USA        Michael Bay    Actio
## 7 3.20e+06      ContentFilm          USA        Wayne Kramer    Crim
## 8 0.00e+00      Orion Pictures          USA        Jim Kouf    Actio
## 9 1.05e+08      Warner Bros.          USA        Roland Emmerich    Actio
## 10 1.70e+06    Magnolia Pictures          USA        Steven Soderbergh    Dram
## 11 1.50e+07      Fox Atomic          USA        Martin Weisz    Actio
## 12 1.60e+07      DC Entertainment          USA        Kenneth Johnson    Actio
## 13 0.00e+00    Colorado Film Production    Italy    Gabriele Salvatores    Crim
## 14 0.00e+00      Caravan Pictures          USA        Martha Coolidge    Comed
## 15 2.00e+07    Magnolia Pictures          UK        Michael Radford    Crim
## # ... with 14 more variables: gross <dbl>, name <fctr>, rating <fctr>,
## #   released <fctr>, runtime <int>, score <dbl>, star <fctr>, votes <int>,
## #   writer <fctr>, year <int>, profit <dbl>, gross_in_mil <dbl>,
## #   budget_in_mil <dbl>, profit_in_mil <dbl>
```

```
# randomly sample a fraction of rows, with replacement
movies %>% sample_frac(0.01, replace=TRUE)
```

```
## # A tibble: 68 x 19
##   budget company country director genre
##   <dbl>   <fctr>   <fctr>   <fctr>   <fctr>
## 1 2.5e+07 Exodus Film Group USA Tony Leondis Animation
## 2 0.0e+00 Chroma III Productions USA Mark Griffiths Comedy
## 3 0.0e+00 Marble Hall USA James Glickenhau Action
## 4 1.2e+05 Chango Productions USA Bryan Johnson Crime
## 5 1.0e+08 Beacon Communications USA Peter Hyams Action
## 6 7.4e+07 Illumination Entertainment USA Kyle Balda Animation
## 7 1.1e+08 Columbia Pictures USA Morten Tyldum Adventure
## 8 0.0e+00 Fox 2000 Pictures USA Adam Collis Comedy
## 9 1.8e+07 Morgan Creek Productions USA John Warren Comedy
## 10 4.9e+07 Carolco Pictures France Paul Verhoeven Drama
## # ... with 58 more rows, and 14 more variables: gross <dbl>, name <fctr>,
## # rating <fctr>, released <fctr>, runtime <int>, score <dbl>, star <fctr>,
## # votes <int>, writer <fctr>, year <int>, profit <dbl>, gross_in_mil <dbl>,
## # budget_in_mil <dbl>, profit_in_mil <dbl>
```