

Lecture 6: Hypothesis testing and Linear Regression

October 19, 2017

Contents

- Hypothesis testing
- Linear Regression
- Extra: Lasso Regression

Useful Books

- “An introduction to Statistical Learning” [ESL] by James, Witten, Hastie and Tibshirani
- “Elements of statistical learning” [ISL] by Hastie, Tibshirani and Friedman
- “Introduction to Linear Regression Analysis” by Montgomery, Peck, Vinning

Hypothesis testing

Hypothesis testing can answer questions:

- **Is the measured quantity equal to/higher/lower than a given threshold?**
e.g. is the number of faulty items in an order statistically higher than the one guaranteed by a manufacturer?
- **Is there a difference between two groups or observations?** e.g. Do treated patient have a higher survival rate than the untreated ones?
- **Is the level of one quantity related to the value of the other quantity?**
e.g. Is hyperactivity related to eating sugar? Is lung cancer related to smoking?

Everyday life examples ([top answer from Quora](#)):

- Test whether route A or route B is the faster way to get from your home to your school.
- Test whether acetaminophen (Tylenol) or ibuprofen (Advil) helps faster with your headaches.
- If you are 21 or older, test whether Tequila or beer gives you worse hangover.
- Test if you run faster in the morning compared to the afternoon.
- Test if your weight is lower in the morning compared to the afternoon.

To perform a hypothesis test you need to:


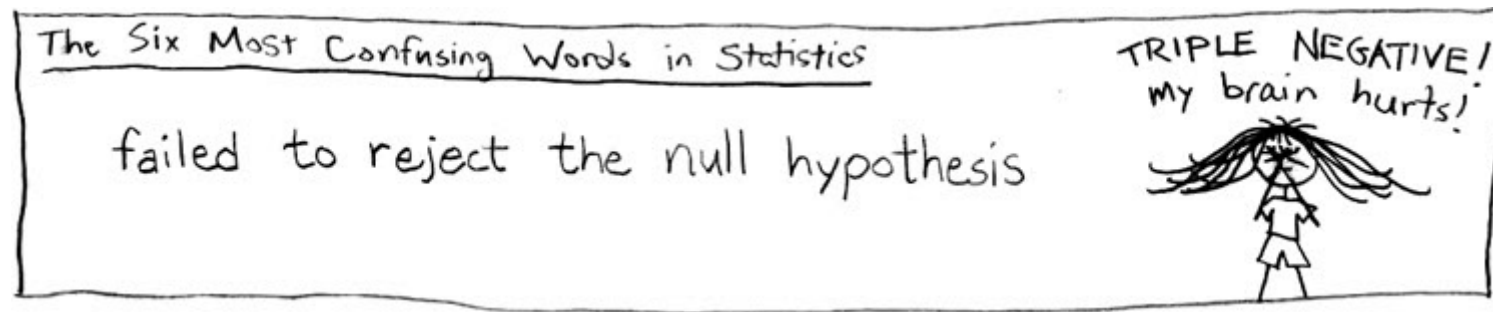
1. Define the null and alternative hypotheses.
2. Choose level of significance α .
3. Pick and compute test statistics.
4. Compute the p-value.
5. Check whether to reject the null hypothesis by comparing p-value to α .
6. Draw conclusion from the test.

Null and alternative hypotheses

The null hypothesis (H_0): A statement assumed to be true unless it can be shown to be incorrect beyond a reasonable doubt. This is something one usually attempts to disprove or discredit.

The alternate hypothesis (H_1): A claim that is contradictory to H_0 and what we conclude when we reject H_0 .

H_0 and H_1 are on purpose set up to be contradictory, so that one can collect and examine data to decide if there is enough evidence to reject the null hypothesis or not.



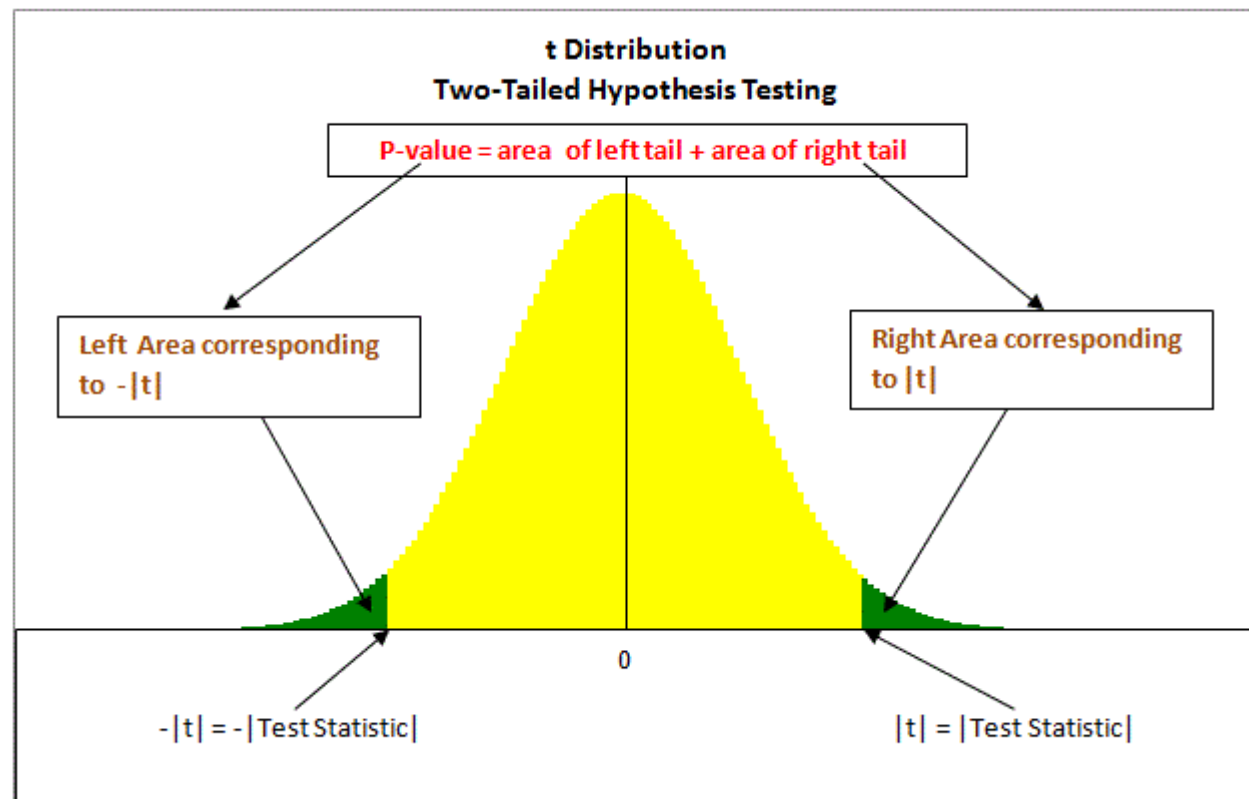
Student's t-test

- William Gosset (1908), a chemist at **the Guinness brewery**.
- Published in Biometrika under a **pseudonym Student**.
- Used to select best yielding varieties of barley.
- Now one of the standard/traditional methods for hypothesis testing.

Distribution of the statistic

p-value is the probability of an observed (or more extreme) result assuming that the null hypothesis is true, i.e. $P[\text{observations} | H_0]$. Note that:

$$P[\text{observations} | \text{hypothesis}] \neq P[\text{hypothesis} | \text{observations}]$$



This is the reason why p-values should NOT be used as a “ranking”/“scoring” system for your hypotheses. You should only use it to potentially reject your hypothesis. Null hypothesis cannot be proven true, you can only fail to reject it.

p-value

- p-value is the **probability of obtaining a result equal to or “more extreme” than what was actually observed, when the null hypothesis is true.**
- A small p-value (typically ≤ 0.05) indicates strong evidence against the null hypothesis, so you reject the null hypothesis.
- A large p-value (> 0.05) indicates weak evidence against the null hypothesis, so you do NOT to reject the null hypothesis.

Dataset

- A built-in dataset, `mtcars`, that comes from a 1974 issue of Motor Trends magazine.

```
data("mtcars")  
head(mtcars)
```

```
##           mpg  cyl  disp  hp  drat    wt    qsec  vs  am  gear  carb  
## Mazda RX4      21.0    6  160  110  3.90  2.620  16.46  0   1     4     4  
## Mazda RX4 Wag  21.0    6  160  110  3.90  2.875  17.02  0   1     4     4  
## Datsun 710      22.8    4  108   93  3.85  2.320  18.61  1   1     4     1  
## Hornet 4 Drive  21.4    6  258  110  3.08  3.215  19.44  1   0     3     1  
## Hornet Sportabout 18.7    8  360  175  3.15  3.440  17.02  0   0     3     2  
## Valiant        18.1    6  225  105  2.76  3.460  20.22  1   0     3     1
```

- rows correspond to car models,
- column are car attributes: miles per gallon, number of cylinders, displacement, transmission etc.

Testing mpg equal to a value


Is the mean fuel efficiency (mpg) in the cars in `mtcars` statistically equal to 25?

Test the null hypothesis:

$$H_0 : \mu = 25$$

$$H_a : \mu \neq 25$$

where μ is the mean mpg of cars in the dataset



```
mtcars$mpg
```

```
## [1] 21.0 21.0 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 17.8 16.4 17.3 15.2
## [15] 10.4 10.4 14.7 32.4 30.4 33.9 21.5 15.5 15.2 13.3 19.2 27.3 26.0 30.4
## [29] 15.8 19.7 15.0 21.4
```

```
tt <- t.test(x = mtcars$mpg, mu = 25, alternative = "two.sided")
tt
```

```
##
## One Sample t-test
##
## data: mtcars$mpg
## t = -4.6079, df = 31, p-value = 6.587e-05
## alternative hypothesis: true mean is not equal to 25
## 95 percent confidence interval:
## 17.91768 22.26357
## sample estimates:
## mean of x
## 20.09062
```

```
names(tt)
```

```
## [1] "statistic" "parameter" "p.value" "conf.int" "estimate"  
## [6] "null.value" "alternative" "method" "data.name"
```

```
# The p-value:  
tt$p.value
```

```
## [1] 6.586738e-05
```

```
# The 95% confidence interval for the mean:  
tt$conf.int
```

```
## [1] 17.91768 22.26357  
## attr(,"conf.level")  
## [1] 0.95
```


Testing mpg smaller than a value


Is the mean fuel efficiency (mpg) in the cars in `mtcars` statistically less than 25?

Test the null hypothesis:

$$H_0 : \mu = 25$$

$$H_a : \mu < 25$$

where μ is the mean mpg of cars in the dataset



```
tt <- t.test(x = mtcars$mpg, mu = 25, alternative = "less")
tt
```

```
##
## One Sample t-test
##
## data: mtcars$mpg
## t = -4.6079, df = 31, p-value = 3.293e-05
## alternative hypothesis: true mean is less than 25
## 95 percent confidence interval:
##      -Inf 21.89707
## sample estimates:
## mean of x
## 20.09062
```

Testing difference between groups

Is the fuel efficiency (mpg) the same for the cars with automatic and manual transmission?

Test the null hypothesis:

$$H_0 : \mu_a = \mu_m$$

$$H_a : \mu_a \neq \mu_m$$

where μ_a mean mpg of automatic cars and μ_m is the mean mpg of manual cars.


Convert the column am (transmission) to a factor:

```
# convert am (0 = automatic, 1 = manual) column to a factor
mtcars <- mtcars %>%
  mutate(am = factor(am, levels = c(0, 1),
                      labels = c("automatic", "manual")))
head(mtcars)
```

```
##      mpg  cyl  disp  hp  drat    wt  qsec  vs          am  gear  carb
##  1  21.0    6   160  110  3.90  2.620  16.46  0      manual    4     4
##  2  21.0    6   160  110  3.90  2.875  17.02  0      manual    4     4
##  3  22.8    4   108   93  3.85  2.320  18.61  1      manual    4     1
##  4  21.4    6   258  110  3.08  3.215  19.44  1  automatic    3     1
##  5  18.7    8   360  175  3.15  3.440  17.02  0  automatic    3     2
##  6  18.1    6   225  105  2.76  3.460  20.22  1  automatic    3     1
```

First, visualize the data

```
library(ggplot2)
ggplot(mtcars, aes(x = am, y = mpg)) + geom_boxplot() +
  xlab("Transmission") + ylab("Fuel efficiency") +
  geom_jitter(width = 0.2, height = 0)
```



The R implementation of the Student's t-test is `t.test()` function:

```
tt <- t.test(x = mtcars$mpg[mtcars$am == "automatic"],  
             y = mtcars$mpg[mtcars$am == "manual"])  
# or  
(tt <- t.test(formula = mpg ~ am, data = mtcars))
```

```
##  
## Welch Two Sample t-test  
##  
## data: mpg by am  
## t = -3.7671, df = 18.332, p-value = 0.001374  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
## -11.280194 -3.209684  
## sample estimates:  
## mean in group automatic mean in group manual  
## 17.14737 24.39231
```

- A tilde symbol, `~`, means “explained by”.

Linear Regression

Linear Regression

- Regression is a supervised learning method, whose goal is inferring the relationship between input data, x , and a **continuous** response variable, y .
- **Linear regression** is a type of regression where **y is modeled as a linear function of x** .
- **Simple linear regression** predicts the output y from a single predictor x .

$$y = \beta_0 + \beta_1 x + \epsilon$$

- **Multiple linear regression** assumes y relies on many covariates:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_p x_p + \epsilon = \vec{\beta} \vec{x} + \epsilon$$

- here ϵ denotes a random noise term with zero mean.

Objective function

Linear regression seeks a solution $\hat{y} = \hat{\beta} \cdot \vec{x}$ that **minimizes the difference between the true outcome y and the prediction \hat{y}** , in terms of the residual sum squares (RSS).

$$\arg \min_{\hat{\beta}} \sum_i \left(y^{(i)} - \hat{\beta} x^{(i)} \right)^2$$

Simple Linear Regression

- Predict the mileage per gallon using the weight of the car.
- In R the linear models can be fit with a `lm()` function.

```
fit <- lm(mpg ~ wt, mtcars)
fit
```

```
##
## Call:
## lm(formula = mpg ~ wt, data = mtcars)
##
## Coefficients:
## (Intercept)          wt
##      37.285      -5.344
```

- Same '`~`' formula notation as for the `t.test` function.

We can check the details on the fitted model by calling:

```
summary(fit)
```

```
##  
## Call:  
## lm(formula = mpg ~ wt, data = mtcars)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -4.5432 -2.3647 -0.1252  1.4096  6.8727   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)  37.2851     1.8776   19.858  < 2e-16 ***  
## wt          -5.3445     0.5591   -9.559  1.29e-10 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 3.046 on 30 degrees of freedom  
## Multiple R-squared:  0.7528, Adjusted R-squared:  0.7446   
## F-statistic: 91.38 on 1 and 30 DF,  p-value: 1.294e-10
```

The coefficients (β) of the model:

```
(beta <- coef(summary(fit)))
```

```
##           Estimate Std. Error   t value    Pr(>|t|)
## (Intercept) 37.285126   1.877627 19.857575 8.241799e-19
## wt         -5.344472   0.559101 -9.559044 1.293959e-10
```

\hat{y} = predicted mpg values for existing observations (cars):

```
predict(fit)[1:15]
```

```
##           Mazda RX4           Mazda RX4 Wag           Datsun 710
##           23.28261           21.91977           24.88595
##           Hornet 4 Drive   Hornet Sportabout           Valiant
##           20.10265           18.90014           18.79325
##           Duster 360           Merc 240D           Merc 230
##           18.20536           20.23626           20.45004
##           Merc 280           Merc 280C           Merc 450SE
##           18.90014           18.90014           15.53313
##           Merc 450SL           Merc 450SLC Cadillac Fleetwood
##           17.35025           17.08302           9.22665
```

To predict the mpg for **new observations**, e.g. a new car with weight $wt = 3.1$ we can do a manual computation using estimated coefficients:

```
beta[, 1]
```

```
## (Intercept)      wt  
##  37.285126    -5.344472
```

```
# beta0 + beta1 * wt  
beta[1, 1] + beta[2, 1]* 3.14
```

```
## [1] 20.50349
```

Predictions can be more efficiently computed using `predict()` function:


```
# create a data frame with new weights:  
newcars <- data.frame(wt = c(2, 2.1, 3.14, 4.1, 4.3))  
predict(fit, newcars)
```

```
##      1      2      3      4      5  
## 26.59618 26.06174 20.50349 15.37279 14.30390
```

```
# Note the same prediction for `wt = 3.14` as computed previously.
```

ggplot2 with `geom_smooth()` can plot the data and the fitted lm model


```
ggplot(mtcars, aes(wt, mpg)) + geom_point() + geom_smooth(method="lm")
```



Multiple Linear Regression

We might like to predict mpg using weight, displacement and the number of cylinders in the car.

```
ggplot(mtcars, aes(x=wt, y=mpg, col=cyl, size=disp)) + geom_point()
```



```
mfit <- lm(mpg ~ wt + disp + cyl, data = mtcars)
```

```
# Summarize the results  
summary(mfit)
```

```
##  
## Call:  
## lm(formula = mpg ~ wt + disp + cyl, data = mtcars)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -4.4035 -1.4028 -0.4955  1.3387  6.0722   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)  41.107678    2.842426  14.462 1.62e-14 ***  
## wt          -3.635677    1.040138   -3.495  0.00160 **  
## disp         0.007473    0.011845    0.631  0.53322   
## cyl         -1.784944    0.607110   -2.940  0.00651 **  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 2.595 on 28 degrees of freedom  
## Multiple R-squared:  0.8326, Adjusted R-squared:  0.8147   
## F-statistic: 46.42 on 3 and 28 DF,  p-value: 5.399e-11
```


To predict mpg for new cars, you must first create a data frame describing the attributes of the new cars:

```
(newcars <- data.frame(wt = c(2, 2.1, 3.14, 4.1, 4.3),  
                        disp = c(100, 200, 500, 300, 210),  
                        cyl = c(6, 6, 4, 6, 8)))
```

```
##      wt  disp cyl  
## 1 2.00  100   6  
## 2 2.10  200   6  
## 3 3.14  500   4  
## 4 4.10  300   6  
## 5 4.30  210   8
```

Then you can compute the predicted mpg

```
predict(mfit, newcars)
```

```
##      1      2      3      4      5  
## 23.87395 24.25768 26.28834 17.73362 12.76403
```

Interaction terms

- An interaction occurs when an independent variable has a different effect on the outcome depending on the values of another independent variable.
- For example, one variable, x_1 might have a different effect on y within different categories or groups, given by variable x_2 .
- If you are not familiar with the concept of the interaction terms, read [this](#).

Models with **interaction effects** can be specified with '*':

```
mfit_inter <- lm(mpg ~ am * wt, mtcars)
summary(mfit_inter)
```

```
##
## Call:
## lm(formula = mpg ~ am * wt, data = mtcars)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.6004  -1.5446  -0.5325   0.9012   6.0909
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   31.4161     3.0201  10.402 4.00e-11 ***
## ammanual      14.8784     4.2640   3.489 0.00162 **
## wt            -3.7859     0.7856  -4.819 4.55e-05 ***
## ammanual:wt   -5.2984     1.4447  -3.667 0.00102 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.591 on 28 degrees of freedom
## Multiple R-squared:  0.833, Adjusted R-squared:  0.8151
## F-statistic: 46.57 on 3 and 28 DF, p-value: 5.209e-11
```

Note that '*' generates the interaction terms:

```
mfit_inter <- lm(mpg ~ am * wt, mtcars)
names(coefficients(mfit_inter))
```

```
## [1] "(Intercept)" "ammanual"      "wt"            "ammanual:wt"
```

You can also specify explicitly which terms you want:

```
mfit_iter2 <- lm(mpg ~ 1 + am + wt + am:wt, mtcars)
summary(mfit_iter2)
```

```
##
## Call:
## lm(formula = mpg ~ 1 + am + wt + am:wt, data = mtcars)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.6004  -1.5446  -0.5325   0.9012   6.0909
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   31.4161     3.0201  10.402 4.00e-11 ***
## ammanual      14.8784     4.2640   3.489 0.00162 **
## wt            -3.7859     0.7856  -4.819 4.55e-05 ***
## ammanual:wt   -5.2984     1.4447  -3.667 0.00102 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.591 on 28 degrees of freedom
## Multiple R-squared:  0.833, Adjusted R-squared:  0.8151
## F-statistic: 46.57 on 3 and 28 DF, p-value: 5.209e-11
```

Extra: Lasso Regression

Choosing a model

- Modern datasets often have “too” many variables, e.g. predict the risk of a disease from the single nucleotide polymorphisms (SNPs) data.
- **Issue:** $n \ll p$ i.e. no. of predictors is much larger than than the no. of observations.
- **Lasso regression** is especially useful for problems, where

the number of available covariates is extremely large, but only a handful of them are relevant for the prediction of the outcome.

Lasso Regression

- Lasso regression is simply regression with L_1 penalty.
- That is, it solves the problem:

$$\hat{\beta}^* = \arg \min_{\hat{\beta}} \sum_i (y^{(i)} - \hat{\beta}x^{(i)})^2 + \lambda \|\hat{\beta}\|_1$$

- It turns out that the L_1 norm $\|\vec{x}\|_1 = \sum_j |x_j|$ **promotes sparsity**.
- The solution, $\hat{\beta}^*$, usually has only a small number of non-zero coefficients.
- The number of non-zero coefficients depends on the choice of the tuning parameter, λ . The higher the λ the fewer non-zero coefficients.

glmnet

- Lasso regression is implemented in an R package glmnet.
- An introductory tutorial to the package can be found [here](#).

```
# install.packages("glmnet")  
library(glmnet)
```

- We go back to `mtcars` datasets and use Lasso regression to predict the mpg using all variables.
- Lasso will pick a subset of predictors (the ones with non-zero coefficients) that best predict the mpg.

```
head(mtcars)
```


```
##      mpg  cyl  disp  hp  drat    wt    qsec  vs          am  gear  carb
## 1  21.0    6   160   110  3.90  2.620  16.46  0     manual     4     4
## 2  21.0    6   160   110  3.90  2.875  17.02  0     manual     4     4
## 3  22.8    4   108    93  3.85  2.320  18.61  1     manual     4     1
## 4  21.4    6   258   110  3.08  3.215  19.44  1  automatic     3     1
## 5  18.7    8   360   175  3.15  3.440  17.02  0  automatic     3     2
## 6  18.1    6   225   105  2.76  3.460  20.22  1  automatic     3     1
```

```
y <- mtcars[, 1] # mileage per gallon
x <- mtcars[, -1] # all other variables treated as predictors
x <- data.matrix(x, "matrix") # converts to NUMERIC matrix
# Choose a training set
set.seed(123)
trainIdx <- sample(1:nrow(mtcars), round(0.7 * nrow(mtcars)))
fit <- glmnet(x[trainIdx, ], y[trainIdx])
names(fit)
```

```
##      [1] "a0"      "beta"    "df"      "dim"      "lambda"
##      [6] "dev.ratio" "nulldev" "npasses" "jerr"     "offset"
##     [11] "call"    "nobs"
```

- `glmnet ()` compute the Lasso regression for a sequence of different tuning parameters, λ .
- Each row of `print (fit)` corresponds to a particular λ in the sequence.
- column `Df` denotes the number of non-zero coefficients (degrees of freedom),
- `%Dev` is the percentage variance explained,
- `Lambda` is the value of the currently chosen tuning parameter.

```
# label = TRUE makes the plot annotate the curves with the corresponding coefficient  
plot(fit, label = TRUE)
```



- the y-axis corresponds the value of the coefficients.
- the x-axis is denoted “ L_1 norm” but is scaled to indicate the number of non-zero coefficients (the effective degrees of freedom).

- Each curve corresponds to a single variable, and shows the value of the coefficient as the tuning parameter varies.
- $\|\hat{\beta}\|_{L_1}$ increases and λ decreases from left to right.
- When λ is small (right) there are more non-zero coefficients.

The computed Lasso coefficient for a particular choice of λ can be printed using

```
# Lambda = 1
coef(fit, s = 1)
```

```
## 11 x 1 sparse Matrix of class "dgCMatrix"
##               1
## (Intercept) 34.877093111
## cyl        -0.867649618
## disp         .
## hp         -0.005778702
## drat         .
## wt         -2.757808266
## qsec         .
## vs           .
## am           .
## gear         .
## carb         .
```

- Like for `lm()`, we can use a function `predict()` to predict the mpg for the training or the test data.
- However, we need specify the value of λ using the argument `s`.

```
# Predict for the test set:  
predict(fit, newx = x[-trainIdx, ], s = c(0.5, 1.5, 2))
```

```
##           1           2           3  
## [1, ] 25.36098 23.87240 23.22262  
## [2, ] 19.82245 19.42427 19.41920  
## [3, ] 16.19324 17.27111 17.74858  
## [4, ] 22.62471 21.86937 21.50396  
## [5, ] 15.20595 16.16123 16.71324  
## [6, ] 11.25687 13.28117 14.26985  
## [7, ] 10.81730 13.01570 14.07314  
## [8, ] 25.88928 24.20103 23.47110  
## [9, ] 27.01880 25.08206 24.22690  
## [10,] 24.89106 23.51713 22.92237
```

Each of the columns corresponds to a choice of λ .

Choosing λ


- To choose λ can use **cross-validation**.
- Use `cv.glmnet()` function to perform a k-fold cross validation.

In k-fold cross-validation, the original sample is randomly partitioned into k equal sized subsamples. Of the k subsamples, a single subsample is retained as the validation data for testing the model, and the remaining k – 1 subsamples are used as training data. ¹


```

set.seed(1)
# `n folds` argument sets the number of folds (k).
cvfit <- cv.glmnet(x[trainIdx, ], y[trainIdx], n folds = 5)
plot(cvfit)

```



- The **red dots** are the average MSE over the k-folds.
- The two chosen λ values are the one with MSE_{min} and one with $MSE_{min} + sd_{min}$

λ with minimum MSE:

```
cvfit$lambda.min
```

```
## [1] 0.2171905
```

The biggest λ such that the MSE is within one standard error of the minimum MSE:

```
cvfit$lambda.1se
```

```
## [1] 0.6632685
```

Exercise I

Exercise I

In this exercise you will perform Lasso regression yourself. We will use the Boston dataset from the MASS package. The dataset contains information on the Boston suburbs housing market collected by David Harrison in 1978.

We will try to predict the median value of homes in the region based on its attributes recorded in other variables.

First install the package:

```
# install.packages("MASS")  
library(MASS)
```

```
head(Boston, 3)
```

```
##      crim zn  indus chas   nox   rm  age   dis rad tax ptratio  black
## 1 0.00632 18   2.31    0 0.538 6.575 65.2 4.0900    1 296    15.3 396.90
## 2 0.02731  0   7.07    0 0.469 6.421 78.9 4.9671    2 242    17.8 396.90
## 3 0.02729  0   7.07    0 0.469 7.185 61.1 4.9671    2 242    17.8 392.83
##    lstat medv
## 1  4.98 24.0
## 2  9.14 21.6
## 3  4.03 34.7
```

```
str(Boston)
```

```
## 'data.frame':    506 obs. of  14 variables:
## $ crim      : num  0.00632 0.02731 0.02729 0.03237 0.06905 ...
## $ zn        : num  18  0  0  0  0  0 12.5 12.5 12.5 12.5 ...
## $ indus     : num  2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 7.87 ...
## $ chas      : int   0  0  0  0  0  0  0  0  0  0 ...
## $ nox       : num  0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524 0.524 ...
## $ rm        : num  6.58 6.42 7.18 7 7.15 ...
## $ age       : num  65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
## $ dis       : num  4.09 4.97 4.97 6.06 6.06 ...
## $ rad       : int   1  2  2  3  3  3  5  5  5  5 ...
## $ tax       : num  296 242 242 222 222 222 311 311 311 311 ...
## $ ptratio   : num  15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 15.2 15.2 ...
## $ black     : num  397 397 393 395 397 ...
## $ lstat     : num  4.98 9.14 4.03 2.94 5.33 ...
## $ medv      : num  24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
```

Split the data to training and testing subsets.

```
set.seed(123)
trainIdx <- sample(1:nrow(Boston), round(0.7 * nrow(Boston)))
boston.test <- Boston[-trainIdx, "medv"]
```

Perform a Lasso regression with `glmnet`. Steps:

1. Extract the input and output data from the `Boston` data frame and convert them if necessary to a correct format.
2. Use cross-validation to select the value for λ .
3. Inspect computed coefficients for `lambda.min`.
4. Compute the predictions for the test dataset the two choices of the tuning parameter, `lambda.min` and `lambda.1se`. Evaluate the MSE for each.

```
# (... ?)
```

1. [https://en.wikipedia.org/wiki/Cross-validation_\(statistics\)#k-fold_cross-validation](https://en.wikipedia.org/wiki/Cross-validation_(statistics)#k-fold_cross-validation)↩