

CME/STATS 195

Lecture 3: Importing and manipulating data

Lan Huong Nguyen

October 4, 2018

Contents

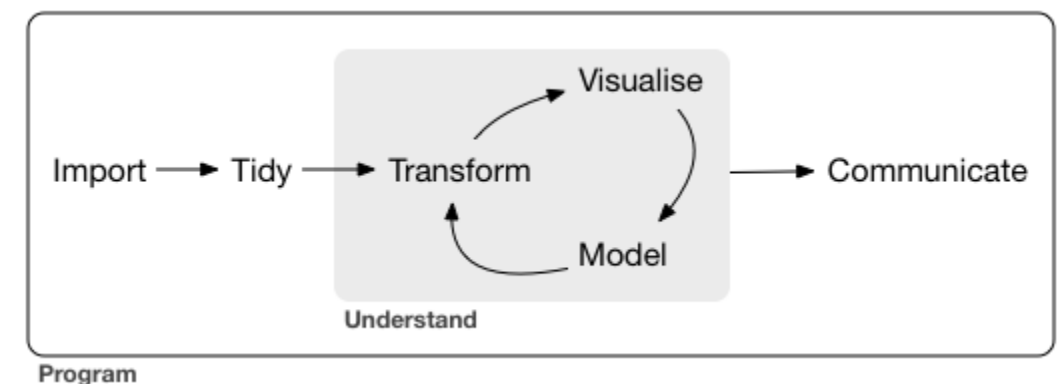
- Data science workflow
- Importing data
- Tidying and manipulating data
- Chaining operations

Data Science

Data Science Workflow

Data science is an exciting discipline that allows you to turn raw data into understanding, insight, and knowledge. ¹

1. Import
2. Wrangle (tidy & transform)
3. Visualize
4. Model
5. Communicate



tidyverse

The tidyverse is an opinionated collection of R packages designed for data science. All packages share an underlying design philosophy, grammar, and data structures. ²

tidyverse includes packages for **importing, wrangling, exploring and modeling data.**

The system is intended to make data scientists more productive. To use tidyverse do the following:

```
# Install the package
install.packages("tidyverse")
# Load it into memory
library("tidyverse")
```



The `tibble` package

The `tibble` package is part of the core `tidyverse`.

Tibbles are a modern take on data frames. They keep the features that have stood the test of time, and drop the features that used to be convenient but are now frustrating.



`tibbles` are data frames, tweaked to make life a little easier. Unlike regular `data.frames` they:

- never change the type of the inputs (e.g. do not convert strings to factors!)
- never changes the names of variables
- never creates `row.names()`
- only recycles inputs of length 1

Using `tibbles`

To use functions from `tibble` and other tidyverse packages:

```
# load it into memory  
library(tidyverse)
```

Printing `tibble` is much nicer, and always fits into your window:

```
# e.g. a built-in dataset 'diamonds' is a tibble:  
class(diamonds)
```

```
## [1] "tbl_df"      "tbl"        "data.frame"
```

```
diamonds
```

```
## # A tibble: 53,940 x 10  
##   carat cut      color clarity depth table price      x      y      z  
##   <dbl> <ord>    <ord> <ord>    <dbl> <dbl> <int> <dbl> <dbl> <dbl>  
## 1 0.23 Ideal     E      SI2      61.5    55   326  3.95  3.98  2.43  
## 2 0.21 Premium  E      SI1      59.8    61   326  3.89  3.84  2.31  
## 3 0.23 Good      E      VS1      56.9    65   327  4.05  4.07  2.31  
## 4 0.290 Premium  I      VS2      62.4    58   334  4.2   4.23  2.63  
## 5 0.31 Good      J      SI2      63.3    58   335  4.34  4.35  2.75  
## 6 0.24 Very Good J      VVS2     62.8    57   336  3.94  3.96  2.48  
## 7 0.24 Very Good I      VVS1     62.3    57   336  3.95  3.98  2.47  
## 8 0.26 Very Good H      SI1      61.9    55   337  4.07  4.11  2.53  
## 9 0.22 Fair      E      VS2      65.1    61   337  3.87  3.78  2.49  
## 10 0.23 Very Good H      VS1      59.4    61   338  4     4.05  2.39  
## # ... with 53,930 more rows
```

Using `tibbles`

Subsetting `tibbles` is stricter than subsetting `data.frames`, and **ALWAYS** returns objects with expected class, i.e. with a single `[` you get back a `tibble`, with double `[]` you get a vector.

```
class(diamonds$carat)
```

```
## [1] "numeric"
```

```
class(diamonds[["carat"]])
```

```
## [1] "numeric"
```

```
class(diamonds[, "carat"])
```

```
## [1] "tbl_df"      "tbl"        "data.frame"
```

```
diamonds.df <- data.frame(diamonds)
class(diamonds.df[["carat"]])
```

```
## [1] "numeric"
```

```
class(diamonds.df[, "carat"])
```



```
## [1] "numeric"
```

More on `tibbles`

You can read more about other `tibble` features by calling on your R console:

```
vignette("tibble")
```

Importing data

Working Directory

- The **current working directory** (cmd) is the location which R is currently pointing to
- Whenever you try to read or save a file without specifying the path explicitly, the cmd will be used by default.
- When are executing code from an R markdown/notebook code chunk, the cmd is **the location of the document**.
- To see the current working directory use `getwd()`:

```
getwd() # with no arguments
```

```
## [1] "/home/lanhuong/MEGA/Teaching/cme195_intro_to_R/cme195.github.io/assets/lectures"
```

- To change the working directory use `setwd(path_name)` with a specified path as na argument:

```
setwd("path/to/directory")
```

Paths and directory names

- R inherits its file and folder **naming conventions from unix**, and uses forward slashes for the directories, e.g. `/home/lan/folder/`
- This is, because backslashes serve a different purpose; they are used as escape characters to isolate special characters and stop them from being immediately interpreted.
- When working with R on **Windows**, you can use either:
`C:/Path/To/A/File` or `C:\\Path\\To\\A\\File`
- Use a “Tab” for autocompletion to find file paths more easily.
- To avoid problems, directory names should NOT contain spaces and special characters.

Importing text data

- **Text Files in a table format** can be read and saved to a selected variable using a `read.table()` function. Use `?read.table` to learn more about the function.
- A common text file format is a **comma delimited text file**, `.csv`. These files use a comma as column separators, e.g:

```
Year,Student,Major
2009, John Doe,Statistics
2009, Bart Simpson, Mathematics I
```

- To read these files use the following command:

```
mydata <- read.table("path/to/filename.csv", header=TRUE, sep = ",")


# read.csv() has convenient argument defaults for '.csv' files
mydata <- read.csv("path/to/filename.csv")
```

- Optionally, use `row.names` or `col.names` arguments to set the row and column names.

The `readr` package

Many R packages provide examples of data. However, sooner or later you will need to work with your own data.

`readr` is for reading rectangular text data into R.



`readr` supports several file formats with seven `read_<...>` functions:

- `read_csv()`: comma-separated (CSV) files
- `read_tsv()`: tab-separated files
- `read_delim()`: general delimited files
- `read_fwf()`: fixed-width files
- `read_table()`: tabular files where columns are separated by white-space
- `read_log()`: web log files

In many cases it just works: supply path to a file and get a tibble back.

Comparison with base R

Why are we learning the readr package?

- it is up to 10x faster
- it produces tibbles instead of data.frames
- better parsing (e.g. does not convert strings to factors)
- more reproducible on different systems
- progress bar for large files

Reading comma-separated files

All `read_<...>()` functions have a similar syntax, so we focus on `read_csv`

```
# Get path to example dataset  
readr_example("mtcars.csv")
```

```
## [1] "/home/lanhuong/R/x86_64-pc-linux-gnu-library/3.4/readr/extdata/mtcars.csv"
```

```
mtcars <- read_csv(readr_example("mtcars.csv"))
```

```
## Parsed with column specification:  
## cols(  
##   mpg = col_double(),  
##   cyl = col_integer(),  
##   disp = col_double(),  
##   hp = col_integer(),  
##   drat = col_double(),  
##   wt = col_double(),  
##   qsec = col_double(),  
##   vs = col_integer(),  
##   am = col_integer(),  
##   gear = col_integer(),  
##   carb = col_integer()  
## )
```

`mtcars` is a dataset on fuel consumption, and other 10 aspects of design and performance (`?mtcars`).

The `read_csv()` function

Also works with inline csv files (useful for experimenting).

```
read_csv(  
  "a,b,c  
1,2,3  
4,5,6"  
)
```

```
## # A tibble: 2 x 3  
##       a     b     c  
##   <int> <int> <int>  
## 1     1     2     3  
## 2     4     5     6
```

```
read_csv(  
  "a,b,c  
1,2,3  
4,5,6",  
  col_names=FALSE  
)
```

```
## # A tibble: 3 x 3  
##      X1     X2     X3  
##   <chr> <chr> <chr>  
## 1 a     b     c  
## 2 1     2     3  
## 3 4     5     6
```

Other useful arguments: `skip_lines`, `symbol` for missing data.

Now you can read most CSV files, also easily adapt to `read_tsv()`, `read_fwf()`. For the others, you need to know how `readr` works inside.

How readr parses data?

```
parse_logical(c("TRUE", "FALSE"))
```

```
## [1] TRUE FALSE
```

```
parse_integer(c("1", "2", "3", "NA"))
```

```
## [1] 1 2 3 NA
```

Parsing vectors:

- `parse_logical()`, `parse_integer()`
- `parse_double()`, `parse_number()`: for numbers from other countries
- `parse_character()`: for character encodings.
- `parse_datetime()`, `parse_date()`, `parse_time()`
- `parse_factor()`

Potential difficulties

Parsing data is not always trivial:

- Numbers are written differently in different parts of the world (“,” vs “.” for separating thousands)
- Numbers are often surrounded by other characters (“\$1000”, “10%”)
- Numbers often contain “grouping” characters (“1,000,000”)
- There are many different ways of writing dates and times
- Times can be in different timezones
- Encodings: special characters in other languages

Locales

A locale specifies common options varying between languages and places

To create a new locale, you use the `locale()` function:

```
locale(  
  date_names = "en",  
  date_format = "%AD",  
  time_format = "%AT",  
  decimal_mark = ".",  
  grouping_mark = ",",  
  tz = "UTC",  
  encoding = "UTF-8",  
  asciify = FALSE)
```

```
## <locale>  
## Numbers: 123,456.78  
## Formats: %AD / %AT  
## Timezone: UTC  
## Encoding: UTF-8  
## <date_names>  
## Days: Sunday (Sun), Monday (Mon), Tuesday (Tue), Wednesday (Wed), Thursday  
##        (Thu), Friday (Fri), Saturday (Sat)  
## Months: January (Jan), February (Feb), March (Mar), April (Apr), May (May),  
##         June (Jun), July (Jul), August (Aug), September (Sep), October  
##         (Oct), November (Nov), December (Dec)  
## AM/PM: AM/PM
```

More on locales can be found in a vignette

```
vignette("locales")
```

Parsing dates

`parse_date()` expects a four digit year, month, day separated by “-” or “/”:

```
parse_date("2010-10-01")
```

```
## [1] "2010-10-01"
```

Example: French format with full name of month:

```
parse_date("1 janvier 2010")
```

```
## Warning: 1 parsing failure.
```

```
## row # A tibble: 1 x 4 col      row    col expected      actual      expected  <int> <int> <chr>      <c
```

```
## [1] NA
```

```
parse_date("1 janvier 2010", format="%d %B %Y", locale=locale("fr"))
```

```
## [1] "2010-01-01"
```

Learn more by typing `?parse_date`

Parsing times

`parse_time()` expects an “hour : minutes” pair (optionally proceeded by “:seconds”, and “am/pm” specifier).

```
parse_time("01:10 am")
```

```
## 01:10:00
```

Parsing dates and times:

```
parse_datetime("2001-10-10 20:10", locale = locale(tz = "Europe/Dublin"))
```

```
## [1] "2001-10-10 20:10:00 IST"
```

For more details, see the book [R for data science](#) or use the documentation.

Parsing numbers

`parse_number()` ignores non-numeric characters before and after.

```
parse_number("20%")
```

```
## [1] 20
```

```
parse_number("$100")
```

```
## [1] 100
```

```
parse_number("cost: $123.45")
```

```
## [1] 123.45
```

Parsing numbers with locales

```
# Separation used in Switzerland  
parse_number("123'456'789", locale = locale(grouping_mark = "'"))
```

```
## [1] 123456789
```


Parsing real numbers

Real numbers using a different decimal mark

```
parse_double("1,23")
```

```
## Warning: 1 parsing failure.
```

```
## row # A tibble: 1 x 4 col      row    col expected      actual expected <int> <int> <chr>
```

```
## [1] NA
```

```
## attr(,"problems")
```

```
## # A tibble: 1 x 4
```

```
##      row    col expected      actual
```

```
##    <int> <int> <chr>      <chr>
```

```
## 1      1    NA no trailing characters ,23
```

```
parse_double("1,23", locale = locale(decimal_mark = ","))
```

```
## [1] 1.23
```

readr's strategy for parsing files

readr uses a heuristic to determine column type, using the first 1000 rows.

You can emulate this process with two functions:

- `guess_parser()`: returns readr's best guess
- `parse_guess()`: uses that guess to parse the column

The heuristic tries a sequence of types, stopping when it finds a match.

If none of these rules apply, then the column will stay as a vector of strings.

```
guess_parser("15:01")
```

```
## [1] "time"
```

```
guess_parser("Oct 10, 2010; 15:01")
```

```
## [1] "character"
```

```
parse_guess("12,352,561")
```

```
## [1] 12352561
```

```
parse_guess(c("TRUE", "FALSE"))
```

```
## [1] TRUE FALSE
```

When the default strategy fails

The default strategy does not always work, e.g. if the first 1000 rows might be a special case. Suppose, your dataset with two columns:

```
# Top 1000 lines are (integer, missing)
readLines(readr_example("challenge.csv"), 10)
```

```
## [1] "x,y"      "404,NA"   "4172,NA"  "3004,NA"  "787,NA"
## [8] "2489,NA"  "1449,NA"  "3665,NA"
```

```
# The remaining are (real number, date)
tail(readLines(readr_example("challenge.csv")), 3000)
```

```
## [1] "0.47193897631950676,2014-08-04" "0.71831864
## [3] "0.26987858884967864,2020-02-04" "0.60823718
```

```
challenge <- read_csv(readr_example("challenge.csv"))
```

```
## Parsed with column specification:
## cols(
##   x = col_integer(),
##   y = col_character()
## )
```

```
## Warning in rbind(names(probs), probs_f): number of columns of result is not a
## multiple of vector length (arg 1)
```

```
## Warning: 1000 parsing failures.
## row # A tibble: 5 x 5 col      row col      expected      actual      file
## ... .....
## See problems(...) for more details.
```

Examining what went wrong

See problems(...) for more details.

`problems(challenge)`

```
## # A tibble: 1,000 x 5
##   row col expected actual file
##   <int> <chr> <chr>      <chr> <chr>
## 1  1001 x      no trailing ch... .238379750... '/home/lanhuong/R/x86_64-pc-linux-g...
## 2  1002 x      no trailing ch... .411679971... '/home/lanhuong/R/x86_64-pc-linux-g...
## 3  1003 x      no trailing ch... .746071676... '/home/lanhuong/R/x86_64-pc-linux-g...
## 4  1004 x      no trailing ch... .723450553... '/home/lanhuong/R/x86_64-pc-linux-g...
## 5  1005 x      no trailing ch... .614524137... '/home/lanhuong/R/x86_64-pc-linux-g...
## 6  1006 x      no trailing ch... .473980569... '/home/lanhuong/R/x86_64-pc-linux-g...
## 7  1007 x      no trailing ch... .578461039... '/home/lanhuong/R/x86_64-pc-linux-g...
## 8  1008 x      no trailing ch... .241593722... '/home/lanhuong/R/x86_64-pc-linux-g...
## 9  1009 x      no trailing ch... .114378662... '/home/lanhuong/R/x86_64-pc-linux-g...
## 10 1010 x      no trailing ch... .298344632... '/home/lanhuong/R/x86_64-pc-linux-g...
## # ... with 990 more rows
```

Fixing the column specifications

Automatic column specifications are:

```
challenge <- read_csv(readr_example("challenge.csv"),
  col_types = cols(x = col_integer(), y = col_character()) )
```

```
## Warning in rbind(names(probs), probs_f): number of columns of result is not a
## multiple of vector length (arg 1)
```

```
## Warning: 1000 parsing failures.
## row # A tibble: 5 x 5 col      row col      expected      actual      file
## ... .....
## See problems(...) for more details.
```

It seems that first column should be a real number:

```
( challenge <- read_csv(readr_example("challenge.csv"),
  col_types = cols(x = col_double(), y = col_character()) ) )
```

```
## # A tibble: 2,000 x 2
##       x y
##   <dbl> <chr>
## 1   404 <NA>
## 2  4172 <NA>
## 3  3004 <NA>
## 4   787 <NA>
## 5    37 <NA>
## 6  2332 <NA>
## 7  2489 <NA>
## 8  1449 <NA>
## 9  3665 <NA>
## 10 3863 <NA>
## # ... with 1,990 more rows
```

Fixing the column specifications

Are we done? Check the “y” column

```
tail(challenge)
```

```
## # A tibble: 6 x 2
##       x y
##   <dbl> <chr>
## 1 0.805 2019-11-21
## 2 0.164 2018-03-29
## 3 0.472 2014-08-04
## 4 0.718 2015-08-16
## 5 0.270 2020-02-04
## 6 0.608 2019-01-06
```

Not yet: dates are stored as strings. To fix this, we use:

```
challenge <- read_csv(readr_example("challenge.csv"),
  col_types = cols(x = col_double(), y = col_date()) )
```

Every `parse_<...>()` function has a corresponding `col_<...>()` function
`col_<...>()` tells readr how to load the data.

Diagnosing problems

Maybe easier to diagnose problems if all columns are read as characters:

```
challenge2 <- read_csv(readr_example("challenge.csv"),  
  col_types = cols(.default = col_character()) )  
head(challenge2, 3)
```

```
## # A tibble: 3 x 2  
##   x      y  
##   <chr> <chr>  
## 1 404   <NA>  
## 2 4172  <NA>  
## 3 3004  <NA>
```

Then use `type_convert()`, which applies the parsing heuristics to the character columns.

```
type_convert(challenge2)
```

```
## Parsed with column specification:  
## cols(  
##   x = col_double(),  
##   y = col_date(format = "")  
## )
```

```
## # A tibble: 2,000 x 2  
##       x y  
##   <dbl> <date>  
## 1   404 NA  
## 2  4172 NA  
## 3  3004 NA
```

```
## 3 3004 NA
## 4 787 NA
## 5 37 NA
## 6 2332 NA
## 7 2489 NA
## 8 1449 NA
## 9 3665 NA
## 10 3863 NA
## # ... with 1,990 more rows
```

Importing other types of data

We will not go into the details in this course. We only list a few other useful packages for importing data.

Rectangular data:

- Package `haven` reads SPSS, Stata, and SAS files.
- Package `readxl` reads excel files (both `.xls` and `.xlsx`).
- Package `DBI`, along with a database specific backend (e.g. `RMySQL`, `RSQLite`, `RPostgreSQL` etc) allows you to run SQL queries against a database and return a data frame.

Hierarchical data:

- `jsonlite` for json (common for browser-server communications)
- `xml2` for XML (common for textual data in web services)

And many more are available.

Exercise 1

- Download “Lec3_Exercises.Rmd” file from the Lectures tab on class website.
- Open the file in RStudio.
- Do Exercise 1.

Tidying and manipulating data

The `dplyr` package

The `dplyr` package is also a part of the core `tidyverse`, which:

- Introduces a **grammar of data manipulation**.
- Gives a **code-efficient** way for data exploration and transformation.
- Is **fast on data frames** (written in C++): has speed of C and ease of R.
- **Intuitive to write and easy to read**, esp. when using the *chaining* syntax.



You should use `dplyr` even as a beginner R user, and [here is why](#).

`dplyr` verbs (functions)

`dplyr` utilities handle the vast majority of your data manipulation needs:

- `filter()` - for picking observations by their values,
- `select()` - for picking variables by their names,
- `arrange()` - for reorder the rows,
- `mutate()` - for creating new variables with functions on existing variables,
- `summarise()` - for collapse many values down to a single summary.

All of the above can be done using **base R functions**, but they would be **less computationally efficient**, and **require writing more lines of (ugly) code**.

The structure of `dplyr` functions

All verbs work similarly:

- The first argument is a tibble (or data frame)
- The subsequent ones describe what to do, using the variable names
- The result is a new tibble

Learn more about `dplyr` from a [tutorial](#) written by its creator, Hadley Wickham

The movie industry dataset

`movies.csv` contains information on last three decades of movies.

The data has been scraped from the IMDb website and can be accessed from a [github repo](https://raw.githubusercontent.com/Juanets/movie-stats/master/movies.csv).

```
url <- "https://raw.githubusercontent.com/Juanets/movie-stats/master/movies.csv"
movies <- read_csv(url)
movies
```

```
## # A tibble: 6,820 x 15
##   budget company country director genre gross name rating released runtime
##   <dbl> <chr>    <chr>    <chr>    <chr>    <dbl> <chr> <chr>    <chr>    <int>
## 1 8.00e6 Columb... USA      Rob Rei... Adve... 5.23e7 Stan... R      1986-08...    89
## 2 6.00e6 Paramo... USA      John Hu... Come... 7.01e7 Ferr... PG-13  1986-06...   103
## 3 1.50e7 Paramo... USA      Tony Sc... Acti... 1.80e8 Top ... PG     1986-05...   110
## 4 1.85e7 Twenti... USA      James C... Acti... 8.52e7 Alie... R      1986-07...   137
## 5 9.00e6 Walt D... USA      Randal ... Adve... 1.86e7 Flig... PG     1986-08...    90
## 6 6.00e6 Hemdale UK        Oliver ... Drama  1.39e8 Plat... R      1987-02...   120
## 7 2.50e7 Henson... UK        Jim Hen... Adve... 1.27e7 Laby... PG     1986-06...   101
## 8 6.00e6 De Lau... USA      David L... Drama  8.55e6 Blue... R      1986-10...   120
## 9 9.00e6 Paramo... USA      Howard ... Come... 4.05e7 Pret... PG-13  1986-02...    96
## 10 1.50e7 SLM Pr... USA      David C... Drama  4.05e7 The ... R      1986-08...    96
## # ... with 6,810 more rows, and 5 more variables: score <dbl>, star <chr>,
## #   votes <int>, writer <chr>, year <int>
```

`filter()`: retain rows matching a criteria

`filter()` allows you to subset observations based on their values.

```
# note: both comma and "&" represent AND condition  
filter(movies, genre == "Comedy", director == "Woody Allen")
```


```
## # A tibble: 27 x 15  
##   budget company country director genre  gross name  rating released runtime  
##   <dbl> <chr>    <chr>    <chr>    <chr>  <dbl> <chr> <chr>    <chr>    <int>  
## 1 6.40e6 Orion  ... USA      Woody A... Come... 4.01e7 Hann... PG-13  1986-03...    107  
## 2 1.60e7 Orion  ... USA      Woody A... Come... 1.48e7 Radi... PG      1987-01...     88  
## 3 1.90e7 Jack R... USA      Woody A... Come... 1.83e7 Crim... PG-13  1989-11...    104  
## 4 1.50e7 Touchs... USA      Woody A... Come... 1.08e7 New ... PG      1989-03...    124  
## 5 1.20e7 Orion  ... USA      Woody A... Come... 7.33e6 Alice PG-13  1991-01...    106  
## 6 1.40e7 Orion  ... USA      Woody A... Come... 2.74e6 Shad... PG-13  1992-03...     85  
## 7 2.00e7 TriSta... USA      Woody A... Come... 1.06e7 Husb... R       1992-09...    103  
## 8 1.35e7 TriSta... USA      Woody A... Come... 1.13e7 Manh... PG      1993-08...    104  
## 9 2.00e7 Miramax USA      Woody A... Come... 1.34e7 Bull... R       1995-02...     98  
## 10 1.50e7 Sweetl... USA      Woody A... Come... 6.70e6 Migh... R       1995-11...     95  
## # ... with 17 more rows, and 5 more variables: score <dbl>, star <chr>,  
## #   votes <int>, writer <chr>, year <int>
```

```
# base R approach would be more wordy:  
movies[movies$genre == "Comedy" & movies$director == "Woody Allen", ]
```

Package `dplyr` executes the filtering and returns a new data frame. **It never modifies the original one.**

Logical operators

Multiple arguments to `filter()` are combined with “and”: all expressions must be true, for a row to be included in the output. For other types of combinations, you need to use Boolean operators yourself: `&` is “and”, `|` is “or”, and `!` is “not”:



Source: [R for data science](#)

```
# Using AND operator
filter(movies, country == "USA", budget > 2.5e8)
# same as filter(movies, country == "USA" & budget > 2.5e8)

# Using OR operator
filter(movies, country == "USA" | budget > 2.5e8)

# Using xor()
filter(movies, xor(score > 9, budget > 2.5e8))
```

```
# you can also use %in% operator
filter(movies, country %in% c("Argentina", "Colombia", "Chile"))
```

```
## # A tibble: 19 x 15
##   budget company country director genre  gross name  rating released runtime
##   <dbl> <chr>    <chr>    <chr>    <chr> <dbl> <chr> <chr>    <chr>    <int>
## 1 0.      Cinequ... Argent... Eliseo ... Drama 7.25e5 Man ... R      1987-03...    105
## 2 0.      "GEA C... Argent... "Mar\xe... Drama 5.21e4 I, t... R      1995-11...    105
## 3 0.      Not sp... Argent... Adolfo ... Drama 1.01e5 A Pl... PG      1994         120
## 4 0.      Aleph ... Argent... Betty K... Drama 1.97e4 De a... R      1996-05...    100
## 5 1.50e6 FX Sou... Argent... "Fabi\x... Crime 1.22e6 Nine... R      2000-08...    114
## 6 0.      Instit... Argent... "Juan J... Come... 6.24e5 Son ... R      2003-01...    123
## 7 0.      4k Fil... Argent... Lucreci... Come... 1.03e5 "La ... NOT R... 2001-04...    103
## 8 0.      FilmFo... Argent... Walter ... Adve... 1.68e7 The ... R      2004-10...    126
## 9 3.00e6 HBO Fi... Colomb... Joshua ... Crime 6.52e6 Mari... R      2004-08...    101
## 10 0.      Cinefa... Argent... Ricardo... Drama 2.80e6 La m... R      2006-04...    100
## 11 0.      Aura F... Argent... "Fabi\x... Crime 5.61e4 The ... UNRAT... 2006-06...    134
## 12 0.      Histor... Argent... "Luc\xe... Drama 4.60e4 XXY   UNRAT... 2008-05...     86
## 13 2.00e6 Tornas... Argent... "Juan J... Drama 2.02e7 The ... R      2010-05...    129
## 14 0.      Partic... Chile    "Pablo ... Drama 2.34e6 No ... R      2012-11...    118
## 15 3.30e6 Corner... Argent... "Dami\x... Come... 3.11e6 Rela... R      2014-12...    122
## 16 2.60e7 Alcon ... Chile    Patrici... Biog... 1.22e7 Los ... PG-13  2015-11...    127
## 17 1.40e6 Buffal... Colomb... Ciro Gu... Adve... 1.33e6 Embr... NOT R... 2015-05...    125
## 18 9.00e6 Fox Se... Chile    "Pablo ... Biog... 1.40e7 Jack... R      2016-12...    100
## 19 0.      AZ Fil... Chile    "Pablo ... Biog... 9.39e5 Neru... R      2017-03...    107
## # ... with 5 more variables: score <dbl>, star <chr>, votes <int>, writer <chr>,
## #   year <int>
```

`select()`: pick columns by name

`select()` lets you choose a subset of variables, specified by name.

Note, there is no need for quotation marks in `dplyr`:

```
# dplyr approach  
select(movies, name, country, year, genre)
```

```
## # A tibble: 6,820 x 4  
##   name                country  year genre  
##   <chr>              <chr>   <int> <chr>  
## 1 Stand by Me        USA      1986 Adventure  
## 2 Ferris Bueller's Day Off USA      1986 Comedy  
## 3 Top Gun            USA      1986 Action  
## 4 Aliens             USA      1986 Action  
## 5 Flight of the Navigator USA      1986 Adventure  
## 6 Platoon            UK       1986 Drama  
## 7 Labyrinth          UK       1986 Adventure  
## 8 Blue Velvet        USA      1986 Drama  
## 9 Pretty in Pink     USA      1986 Comedy  
## 10 The Fly           USA      1986 Drama  
## # ... with 6,810 more rows
```

```
# base R approach would be:  
movies[, c("name", "year", "genre")]
```

```
# use colon to select multiple contiguous columns,
select(movies, name, genre:score)
```

```
## # A tibble: 6,820 x 7
##   name          genre      gross rating released  runtime  score
##   <chr>         <chr>    <dbl> <chr>  <chr>      <int>  <dbl>
## 1 Stand by Me    Adventure 52287414 R      1986-08-22    89    8.1
## 2 Ferris Bueller's Day Off Comedy    70136369 PG-13   1986-06-11   103    7.8
## 3 Top Gun        Action    179800601 PG      1986-05-16   110    6.9
## 4 Aliens          Action    85160248 R       1986-07-18   137    8.4
## 5 Flight of the Navigator Adventure 18564613 PG      1986-08-01    90    6.9
## 6 Platoon         Drama    138530565 R       1987-02-06   120    8.1
## 7 Labyrinth       Adventure 12729917 PG      1986-06-27   101    7.4
## 8 Blue Velvet     Drama      8551228 R       1986-10-23   120    7.8
## 9 Pretty in Pink  Comedy    40471663 PG-13   1986-02-28    96    6.8
## 10 The Fly        Drama     40456565 R       1986-08-15    96    7.5
## # ... with 6,810 more rows
```

```
# To drop columns use a minus, "-"
select(movies, -(star:writer))
```

```
## # A tibble: 6,820 x 12
##   budget company country director genre  gross name rating released runtime
##   <dbl> <chr>    <chr>  <chr>  <chr>  <dbl> <chr> <chr>  <chr>      <int>
## 1 8.00e6 Columb... USA    Rob Rei... Adve... 5.23e7 Stan... R      1986-08...    89
## 2 6.00e6 Paramo... USA    John Hu... Come... 7.01e7 Ferr... PG-13   1986-06...   103
## 3 1.50e7 Paramo... USA    Tony Sc... Acti... 1.80e8 Top ... PG      1986-05...   110
## 4 1.85e7 Twenti... USA    James C... Acti... 8.52e7 Alie... R       1986-07...   137
## 5 9.00e6 Walt D... USA    Randal ... Adve... 1.86e7 Flig... PG      1986-08...    90
## 6 6.00e6 Hemdale UK      Oliver ... Drama 1.39e8 Plat... R       1987-02...   120
## 7 2.50e7 Henson... UK      Jim Hen... Adve... 1.27e7 Laby... PG      1986-06...   101
## 8 6.00e6 De Lau... USA    David L... Drama 8.55e6 Blue... R       1986-10...   120
## 9 9.00e6 Paramo... USA    Howard ... Come... 4.05e7 Pret... PG-13   1986-02...    96
## 10 1.50e7 SLM Pr... USA    David C... Drama 4.05e7 The ... R       1986-08...    96
## # ... with 6,810 more rows, and 2 more variables: score <dbl>, year <int>
```

select() helpers

You can use the following functions to help select the columns:

- `starts_with()`
- `ends_with()`
- `contains()`
- `matches()` (matches a regular expression)
- `num_range("x", 1:4)`: picks variables x1, x2, x3, x4

Examples:

```
select(movies, starts_with("r"))  
select(movies, ends_with("e"))  
select(movies, contains("re"))
```

arrange(): reorder rows

arrange() takes a data frame and a set of column names to order by.
For descending order, use the function desc() around the column name.

```
print(arrange(movies, runtime), n = 4)
```

```
## # A tibble: 6,820 x 15
##   budget company country director genre  gross name  rating released runtime
##   <dbl> <chr>    <chr>    <chr>    <chr> <dbl> <chr> <chr>    <chr>    <int>
## 1 0.      Iwerks... France  Jean-Ja... Adve... 1.51e7 Wing... G      1996-09...    50
## 2 1.25e7 Univer... USA     Don Blu... Anim... 4.81e7 The ... G      1988-11...    69
## 3 6.00e3 Next W... UK      Christo... Crime 4.85e4 Foll... R      1999-11...    69
## 4 0.      Hyperi... USA     Bruce W... Anim... 8.44e6 "B\x... PG-13  1992-07...    70
## # ... with 6,816 more rows, and 5 more variables: score <dbl>, star <chr>,
## #   votes <int>, writer <chr>, year <int>
```

```
# use `desc` for descending
print(arrange(movies, desc(budget)), n = 4)
```

```
## # A tibble: 6,820 x 15
##   budget company country director genre  gross name  rating released runtime
##   <dbl> <chr>    <chr>    <chr>    <chr> <dbl> <chr> <chr>    <chr>    <int>
## 1 3.00e8 Walt D... USA     Gore Ve... Acti... 3.09e8 Pira... PG-13  2007-05...   169
## 2 2.60e8 Walt D... USA     Nathan ... Anim... 2.01e8 Tang... PG      2010-11...   100
## 3 2.58e8 Columb... USA     Sam Rai... Acti... 3.37e8 Spid... PG-13  2007-05...   139
## 4 2.50e8 Warner... UK      David Y... Adve... 3.02e8 Harr... PG      2009-07...   153
## # ... with 6,816 more rows, and 5 more variables: score <dbl>, star <chr>,
## #   votes <int>, writer <chr>, year <int>
```

mutate(): add new variables

mutate() adds new columns that are a function of the existing ones

```
movies <- mutate(movies, profit = gross - budget)
select(movies, name, gross, budget, profit)
```

```
## # A tibble: 6,820 x 4
##   name                gross    budget    profit
##   <chr>              <dbl>    <dbl>    <dbl>
## 1 Stand by Me        52287414  8000000  44287414
## 2 Ferris Bueller's Day Off 70136369  6000000  64136369
## 3 Top Gun            179800601 15000000 164800601
## 4 Aliens              85160248 18500000  6660248
## 5 Flight of the Navigator 18564613  9000000  9564613
## 6 Platoon            138530565  6000000 132530565
## 7 Labyrinth          12729917 25000000 -12270083
## 8 Blue Velvet         8551228  6000000  2551228
## 9 Pretty in Pink      40471663  9000000  31471663
## 10 The Fly            40456565 15000000  25456565
## # ... with 6,810 more rows
```

To discard old variables, use `transmute()` instead of `mutate()`.

```
# base R approach to create a new variable 'profit'
movies$profit <- movies$gross - movies$budget
```



```
# Generating multiple new variables
movies <- mutate(
  movies,
  profit = gross - budget,
  gross_in_mil = gross/10^6,
  budget_in_mil = budget/10^6,
  profit_in_mil = profit/10^6
)
select(movies, name, year, country, contains("_in_mil"), profit)
```

```
## # A tibble: 6,820 x 7
##   name      year country gross_in_mil budget_in_mil profit_in_mil  profit
##   <chr>    <int> <chr>      <dbl>         <dbl>         <dbl>    <dbl>
## 1 Stand by Me  1986 USA         52.3           8         44.3    4.43e7
## 2 Ferris Buel... 1986 USA         70.1           6         64.1    6.41e7
## 3 Top Gun      1986 USA        180.          15        165.    1.65e8
## 4 Aliens       1986 USA         85.2          18.5        66.7    6.67e7
## 5 Flight of t... 1986 USA         18.6           9          9.56    9.56e6
## 6 Platoon      1986 UK         139.           6        133.    1.33e8
## 7 Labyrinth    1986 UK          12.7          25        -12.3   -1.23e7
## 8 Blue Velvet  1986 USA          8.55           6          2.55    2.55e6
## 9 Pretty in P... 1986 USA         40.5           9         31.5    3.15e7
## 10 The Fly     1986 USA         40.5          15         25.5    2.55e7
## # ... with 6,810 more rows
```

Any vectorized function can be used with **mutate()**, including:

- arithmetic operators (+, -, *, /, %, %%),
- logical operators (<, <=, >, >=, ==, !=),
- logarithmic and exponential transformations (log, log10, exp),
- offsets (lead, lag),
- cumulative rolling aggregates (cumsum, cumprod, cummin, cummax),
- ranking (min_rank, percent_rank).

`summarise()`: reduce variables to values

`summarize()` can be used to aggregate data or to compute a summarizing value of interest.

```
summarise(movies,  
  tot_gross_in_bil = sum(gross)/1e9,  
  mean_gross_in_mil = mean(gross)/1e6,  
  mean_profit_in_mil = mean(profit)/1e6  
)
```

```
## # A tibble: 1 x 3  
##   tot_gross_in_bil mean_gross_in_mil mean_profit_in_mil  
##               <dbl>             <dbl>             <dbl>  
## 1             228.              33.5              8.92
```

`summarize()` is more useful on data previously grouped by one or more variables using `group_by()`.

```
by_genre <- group_by(movies, genre)  
summarize(by_genre, tot_gross_in_bil = sum(gross)/1e9,  
  mean_gross_in_mil = mean(gross)/1e6,  
  mean_profit_in_mil = mean(profit)/1e6)
```

Grouping and summarizing

Grouping allows you to compute summaries for each categories separately:

```
by_genre <- group_by(movies, genre)
summarize(
  by_genre,
  tot_gross_in_bil = sum(gross)/1e9,
  mean_gross_in_mil = mean(gross)/1e6,
  mean_profit_in_mil = mean(profit)/1e6
)
```

```
## # A tibble: 17 x 4
##   genre      tot_gross_in_bil mean_gross_in_mil mean_profit_in_mil
##   <chr>          <dbl>          <dbl>          <dbl>
## 1 Action          74.8           56.2           7.30
## 2 Adventure       20.9           53.3          16.0
## 3 Animation       25.3           91.5          27.2
## 4 Biography        8.62           24.0           7.05
## 5 Comedy          53.5           25.7          10.8
## 6 Crime           10.2           19.6           3.30
## 7 Drama           25.2           17.5           4.19
## 8 Family           0.118            8.44          -0.101
## 9 Fantasy           0.645           20.1           4.38
## 10 Horror           7.12           25.7          13.8
## 11 Musical          0.00809           2.02          -0.476
## 12 Mystery           1.38           36.3           9.47
## 13 Romance           0.146            9.72           4.24
## 14 Sci-Fi           0.308           23.7           6.79
## 15 Thriller          0.0996            5.53          -0.356
## 16 War              0.00151            0.755           0.755
## 17 Western           0.0185            9.26           3.26
```

Elementary but useful summary functions

- `min(x), median(x), max(x), quantile(x, p)`
- `n(), n_distinct(), sum(x), mean(x)`
- `sum(x > 10), mean(x > 0)`
- `sd(x), var(x)`

Counting observations

`tally()` function can be used to generate a group frequency table, (number of observations in each category)

```
tally(group_by(movies, genre))
```

```
## # A tibble: 17 x 2
##   genre      n
##   <chr>    <int>
## 1 Action   1331
## 2 Adventure  392
## 3 Animation  277
## 4 Biography  359
## 5 Comedy   2080
## 6 Crime     522
## 7 Drama    1444
## 8 Family     14
## 9 Fantasy    32
## 10 Horror    277
## 11 Musical     4
## 12 Mystery    38
## 13 Romance    15
## 14 Sci-Fi     13
## 15 Thriller   18
## 16 War         2
## 17 Western     2
```

```
tally(group_by(movies, genre, country))
```

```
## # A tibble: 238 x 3
## # Groups:   genre [17]
##   genre country      n
##   <chr> <chr>    <int>
## 1 Action Aruba         1
## 2 Action Australia    12
## 3 Action Austria       1
## 4 Action Belgium       1
## 5 Action Brazil        2
## 6 Action Canada       26
## 7 Action China        13
## 8 Action Czech Republic 1
## 9 Action Denmark       2
## 10 Action France       41
## # ... with 228 more rows
```

Window Functions

- Aggregation functions such as `mean()`, `n()` return 1 value per group.
- **Window functions return multiple values per group**, e.g. `top_n()`, `lead` and `lag` or `cummean`:

```
# rewrite more simply with the `top_n` function
movies2 <- select(movies, name, genre, country, year, budget, gross, profit, rating, score)
top2 <- top_n(group_by(movies2, genre), n = 2, wt = score)
arrange(top2, genre, year, score)
```

```
## # A tibble: 35 x 9
## # Groups:   genre [17]
##   name          genre country    year budget  gross  profit rating  score
##   <chr>         <chr>   <chr>   <int>  <dbl>  <dbl>   <dbl> <chr>   <dbl>
## 1 The Dark Knight Action   USA     2008 1.85e8 5.35e8 3.50e8 PG-13     9
## 2 Inception      Action   USA     2010 1.60e8 2.93e8 1.33e8 PG-13    8.8
## 3 The Lord of the ... Advent... New Zeal... 2001 9.30e7 3.16e8 2.23e8 PG-13    8.8
## 4 The Lord of the ... Advent... USA       2003 9.40e7 3.78e8 2.84e8 PG-13    8.9
## 5 The Lion King   Animat... USA       1994 4.50e7 3.13e8 2.68e8 G        8.5
## 6 Spirited Away   Animat... Japan     2001 1.90e7 1.01e7 -8.94e6 PG       8.6
## 7 Your name       Animat... Japan     2016 0.      5.02e6 5.02e6 PG       8.5
## 8 Schindler's List Biogra... USA       1993 2.20e7 9.61e7 7.41e7 R        8.9
## 9 The Intouchables Biogra... France    2011 0.      1.32e7 1.32e7 R        8.6
## 10 Forrest Gump    Comedy   USA       1994 5.50e7 3.30e8 2.75e8 PG-13    8.8
## # ... with 25 more rows
```

Other useful functions in dplyr

```
# Renaming variables
```

```
print(rename(movies2, gross_revenue = gross), n = 5)
```

```
## # A tibble: 6,820 x 9
##   name          genre country year budget gross_revenue profit rating score
##   <chr>         <chr>   <chr> <int> <dbl>         <dbl>   <dbl> <chr>  <dbl>
## 1 Stand by Me   Advent... USA    1986 8.00e6      52287414 4.43e7 R      8.1
## 2 Ferris Buell... Comedy   USA    1986 6.00e6      70136369 6.41e7 PG-13   7.8
## 3 Top Gun       Action   USA    1986 1.50e7     179800601 1.65e8 PG      6.9
## 4 Aliens        Action   USA    1986 1.85e7      85160248 6.67e7 R      8.4
## 5 Flight of th... Advent... USA    1986 9.00e6     18564613 9.56e6 PG      6.9
## # ... with 6,815 more rows
```

```
# Unique values
```

```
distinct(movies2, rating)
```

```
## # A tibble: 13 x 1
##   rating
##   <chr>
## 1 R
## 2 PG-13
## 3 PG
## 4 UNRATED
## 5 Not specified
## 6 G
## 7 NC-17
## 8 NOT RATED
## 9 TV-PG
## 10 TV-MA
## 11 B
## 12 B15
## 13 TV-14
```

```
# Using multiple variables, returns distinct variables
```

```
distinct(movies2, rating, genre)
```

```
## # A tibble: 83 x 2
##   rating genre
##   <chr>   <chr>
## 1 R      Adventure
## 2 PG-13  Comedy
## 3 PG      Action
## 4 R      Action
## 5 PG      Adventure
## 6 R      Drama
## 7 PG-13  Adventure
## 8 PG-13  Action
## 9 R      Crime
## 10 UNRATED Comedy
## # ... with 73 more rows
```

Sampling observations

```
sample_n(movies, 8) # fixed number of rows, without replacement
```

```
## # A tibble: 8 x 19
##   budget company country director genre  gross name  rating released runtime
##   <dbl> <chr>    <chr>    <chr>    <chr>  <dbl> <chr> <chr>    <chr>    <int>
## 1 0.      Incent... USA      Don Roos Drama  2.06e4 Love... R      2011-03...    119
## 2 3.20e7 David ... USA      Mark Di... Anim... 3.56e6 Cats... G      1997-03...     75
## 3 8.00e6 Avenue... USA      Robert ... Come... 2.17e7 The ... R      1992-05...    124
## 4 1.10e7 Motlys  Norway Joachim... Drama  1.62e5 "M\x... R      2016-07...    109
## 5 0.      Britis... UK       Terence... Biog... 1.61e5 The ... PG     1993-05...     85
## 6 2.10e7 Scott ... USA      Trey Pa... Anim... 5.20e7 Sout... R      1999-06...     81
## 7 0.      Scary ... USA      Jeffrey... Come... 2.75e6 Dead... R      1986-11...     93
## 8 4.50e7 R.P. P... France Roman P... Drama  1.55e7 The ... PG-13  2010-03...    128
## # ... with 9 more variables: score <dbl>, star <chr>, votes <int>, writer <chr>,
## #   year <int>, profit <dbl>, gross_in_mil <dbl>, budget_in_mil <dbl>,
## #   profit_in_mil <dbl>
```

```
sample_frac(movies, 0.005, replace=TRUE) # fraction of rows, with replacement
```

```
## # A tibble: 34 x 19
##   budget company country director genre  gross name  rating released runtime
##   <dbl> <chr>    <chr>    <chr>    <chr>  <dbl> <chr> <chr>    <chr>    <int>
## 1 1.00e7 Avenue... USA      Martin ... Crime  3.84e5 Amer... R      1993-05...    113
## 2 8.00e6 Odeon ... Canada David C... Drama  1.64e6 Spid... R      2002-11...     98
## 3 2.50e7 Twenti... USA      Jim Abr... Acti... 3.89e7 Hot ... PG-13  1993-05...     86
## 4 5.50e7 Paramo... USA      Robert ... Come... 3.30e8 Forr... PG-13  1994-07...    142
## 5 2.50e7 Fox 20... USA      Jonatha... Drama  1.01e7 Brok... PG-13  1999-08...    100
## 6 0.      Rysher... USA      John He... Come... 1.11e7 2 Da... R      1996-09...    104
## 7 1.35e8 DreamW... USA      Kirk De... Anim... 1.87e8 The ... PG     2013-03...     98
## 8 2.50e7 Screen... USA      Dwight ... Acti... 3.22e7 Anac... PG-13  2004-08...     97
## 9 0.      Davros... USA      Euzhan ... Drama  3.77e6 A Dr... R      1989-09...     97
## 10 2.50e5 Paramo... USA      Drake D... Drama  3.41e6 Like... PG-13  2012-01...     86
## # ... with 24 more rows, and 9 more variables: score <dbl>, star <chr>,
## #   votes <int>, writer <chr>, year <int>, profit <dbl>, gross_in_mil <dbl>,
## #   budget_in_mil <dbl>, profit_in_mil <dbl>
```

```
## # budget_in_m11 <dbl>, profit_in_m11 <dbl>
```

Chaining operations

The magrittr package


The magrittr (to be pronounced with a sophisticated french accent) package has two aims: decrease development time and improve readability and maintainability of code



magrittr provides a “pipe”-like operator, %>%:

- The %>% is used **pipe values forward into an expression or function call**.
- In the pipe notation, you use `x %>% f(y)`, rather than `f(x, y)`.
- This is similar to the [Unix pipes](#), |, used to send the output of one program to another program for further processing.

```
columns <- 1:50  
rnorm(500) %>%  
  matrix(ncol = 50) %>%  
  colSums() %>%  
  plot(x = columns)
```



Chaining operations

- Pipe operators used together with `dplyr` functions make a large difference as they semantically change your code in a way that **makes it more intuitive to both read and write**.
- The pipes allow users to chain operators which reflects the **sequential nature of data-processing tasks**.
- **Chaining increases readability** significantly when there are many commands
- `%>%` operator is automatically imported into `dplyr`

1. Find movies from USA produced after 2010.
- (2) Group by genre and compute the group mean gross revenue in million dollars. Then print the genre mean 'gross' revenue (4) arranged in a descending order:

```
# nesting
arrange(
  summarise(
    group_by(
      filter(movies,
              year > 2010, country == "USA"
            ),
      genre
    ),
    mean_gross = mean(gross)/10^6
  ),
  mean_gross
)
```

```
## # A tibble: 13 x 2
##   genre      mean_gross
##   <chr>      <dbl>
## 1 Thriller    0.0165
## 2 Drama      23.3
## 3 Horror     27.7
## 4 Sci-Fi     29.2
## 5 Fantasy    30.7
## 6 Crime      32.1
## 7 Comedy     35.2
## 8 Biography  40.6
## 9 Mystery    49.5
## 10 Romance   62.5
## 11 Adventure  81.2
## 12 Action    97.3
## 13 Animation 152.
```

```
# chaining
movies %>%
  filter(year > 2010, country == "USA") %>%
  group_by(genre) %>%
  summarise(mean_gross = mean(gross)/10^6) %>%
  arrange(mean_gross)
```

```
## # A tibble: 13 x 2
##   genre      mean_gross
##   <chr>      <dbl>
## 1 Thriller    0.0165
## 2 Drama      23.3
## 3 Horror     27.7
## 4 Sci-Fi     29.2
## 5 Fantasy    30.7
## 6 Crime      32.1
## 7 Comedy     35.2
## 8 Biography  40.6
## 9 Mystery    49.5
## 10 Romance   62.5
## 11 Adventure  81.2
## 12 Action    97.3
## 13 Animation 152.
```

Saving the workspace

- You can choose to **save all objects** currently in the workspace (variables, functions, etc.) into a file e.g. `filename.rda`.
- The file `filename.rda` can be loaded next time you work with R.
- You can also save a single object or a subset of specified objects currently in the workspace.

```
# save the workspace to file  
save.image(file = "path/to/filename.rda")  
  
# save specific objects to a file  
save(object_list, file = "path/to/filename.rda")  
  
# save just a single object  
saveRDS(object, file = "path/to/filename.rds")
```

- Saved objects/workspace can be loaded back in a new R session.

```
# load a workspace into the current session  
load("path/to/filename.rda")  
  
# read just the previously saved 1 object  
object <- readRDS("path/to/filename.rds")
```

Exercises 2

- Go to the “Lec3_Exercises.Rmd” file, which can be downloaded from the class website under the Lecture tab.
- Complete Exercise 2.

-
1. R for Data Science↵
 2. Tidyverse website↵